Optimising Binary HVAC Switching in Saudi Residential Buildings with Model Predictive Control and Rule-Constrained PPO

Hamzah Faraj[1*]

[1*]Department of Sciences and Technology, Ranyah University College, Taif University, P.O. Box 11099, Taif, 21944, Saudi Arabia.

Abstract

On/off cooling of Saudi split units is posed as constrained optimisation with hard comfort constraints, an energy cost and a charge per on-command, and evaluated in Riyadh, Jeddah and Taif on held-out weeks against an assumed **1** °C-deadband thermostat shifted to each controller's mean temperature. Tuning its two switch points widens the deadband to **1.32**–**1.36** °C, near the widest a **15**-min step admits; it avoids **22**–**24%** of on-commands with electricity within **±0.5%**, worth **2**–**19** Saudi riyals (SAR) per room per season at the **0.025**–**0.114** SAR per on-command that a wear estimate supports. Its cost saving, **12.4**–**13.6%** at a **0.15** SAR reference charge, falls to **5.7%** against a **1.25** °C comparator (Riyadh). A **24** h model predictive controller (MPC) avoids **30**–**33%** with a perfect in-horizon forecast; with persistence forecasts its margin is not resolved. In Riyadh the tuned thermostat saves **−2.8%** at half and **15.8%** at twice the reference capacity, probably above the capacity acting per cycle, and MPC's margin shrinks from **15.0** to **0.4** points; at a **5**-min step it saves about **25%**, at **30** min **−1.2** to **−1.5%**, where MPC keeps **6.6**–**11.4%**. Proximal policy optimisation inside hard rules reproduces their hold policy in two cities and falls below it in Taif. A clairvoyant schedule and a provable bound place the best saving at **18.0**–**20.0%**, **18.4**–**19.9%** and **16.7**–**21.0%**. In two- and four-room dwellings the tuned thermostat saves **11.7**–**13.6%** (one room: **12.5%**; not resolved). Electricity follows the comfort target: each **0.1** K warmer mean lowers the thermostat's by **0.9**–**2.3%**.

Keywords: Binary HVAC switching, model predictive control, constrained reinforcement learning, Saudi Vision 2030, Green Riyadh, CO_2 reduction

1 Introduction

Buildings consume approximately 40% of global final energy, with HVAC systems accounting for roughly half that figure (Pérez-Lombard et al. 2008). In Saudi Arabia the share is larger: residential buildings alone consume ≈50% of national electricity, and air conditioning represents up to 70% of residential demand (Al-Homoud and Krarti 2021; Krarti et al. 2020). Better operation of residential air conditioning is therefore one route to the efficiency targets of Saudi Vision 2030 (Belaïd and Fikru 2025), whose Saudi Green Initiative commits to cutting emissions by more than 278 Mt a year by 2030 and to net zero by 2060 (Saudi Green Initiative 2021); in the capital, the Green Riyadh programme of more than 7.5 million trees aims to lower the air temperature by 1.5–2 °C (Saudipedia 2019). Section 10.3 separates what switching control and what the comfort target and a cooler city each contribute to those aims.

1.1 Motivation

The prevalent residential control strategy is the automatic thermostat: it switches on when a room exceeds an upper setpoint and off when it returns, with no tariff awareness and no inter-zone coordination. It cycles on a fixed deadband, and every switch-on command adds wear that it does not price. Model predictive control (MPC) plans against a model and a forecast (Drgoňa et al. 2020); reinforcement learning (RL) learns a policy from interaction without an explicit model (Al Sayed et al. 2024; Biemann et al. 2021; Azuatalam et al. 2020), yet pure RL agents can violate comfort constraints during exploration (Xu et al. 2025). Hybrid approaches combine hard safety rules with RL (Xu et al. 2025). The closest prior work is the present author's physics-informed proximal policy optimisation (PPO) framework for Saudi residential split units under the same tariff (Faraj 2026), which enforces comfort through a soft physics penalty and reports no optimality certificate; Section 2.4 sets out what the present study adds to it and what it takes from it.

1.2 Contributions

Six contributions are made. (C1) A constrained binary formulation of residential on/off cooling, over single- and multi-zone topologies, with a hard occupied ceiling and a hard cold rule (a cooling step may never end below the band; free-float is unrestricted), a charge per supervisory on-command and a tariff-weighted energy cost, and three optimality conditions under a limit-average criterion. (C2) A connection to Simple Linear Hybrid Automata (Mousa et al. 2016) which shows that, under constant disturbance and at the mean of its own maximal-width cycle, the hysteresis induced by the hard rules is optimal among policies that act only where the rules leave the action free (Corollary 1), and which names the features that remove that optimality, a diurnal disturbance first among them. (C3) One selection and evaluation protocol applied identically to every admitted arm (E2 and the MANUAL setback reference, which are not comfort-matched, are reported outside it): parameters chosen on tuning weeks under a single admissibility rule with a single fallback, and each held-out week compared with a thermostat shifted to the arm's own clipped mean occupied temperature, with pooled savings and exact paired statistics in pre-declared families. (C4) A controlled comparison of a tuned two-parameter thermostat, rule-constrained PPO, a one-step model shield and receding-horizon (MPC) controllers from 15 min to 48 h, with perfect and persistence forecasts and across thermal capacities. The tuned thermostat's gain is a wider effective deadband. The receding-horizon controller's margin over it is resolved only with a perfect in-horizon forecast, shrinks with thermal capacity and changes with the actuation step. An exact one-step shield in place of the rules' fixed margins matches the tuned thermostat, and PPO inside the rules reproduces the rules' hold policy in Riyadh and Jeddah and falls below it in Taif. (C5) A clairvoyant optimality bracket, from the best attained schedule no warmer than the baseline to a provable relaxation bound, with the bound also reported for every arm at its own mean temperature. (C6) A three-layer report of every result (on-commands avoided, electricity, and money as a range over the supported charge per on-command), reporting that at matched mean temperature the cost difference is a charge on avoided on-commands, which by an identity of the model leaves electricity unchanged, and a Riyadh check of whether the single-zone ordering survives in dwellings of two and four rooms with independent units.

Section 2 surveys the related literature. Section 3 defines the thermal and cost models. Section 4 formalises the constrained optimisation. Section 5 sets out the controllers. Section 6 characterises the three climates. Section 7 states the protocol. Section 8 reports results. Section 9 compares the results with published figures and reports the ablations and sensitivities. Section 10 interprets the findings and states their limitations. Section 11 concludes.

2 Literature Review

The review sets out the Saudi context, the control approaches, binary actuation with multi-zone coupling and tariffs, and then the gap, which has four components: binary actuation, hard comfort enforcement, an optimality reference, and Saudi climate and tariff coverage.

2.1 Saudi Context

Al-Homoud and Krarti (2021) attribute the Kingdom's residential demand to a hot climate, low energy awareness and a historically subsidised tariff. Residential electricity is priced on a step-wise schedule (Nahiduzzaman et al. 2023; Saudi Gazette 2025): the four-tier schedule of 2016–2017 rose from 0.05 to 0.30 SAR/kWh (Nahiduzzaman et al. 2023); the 2018 reform set 0.18 SAR/kWh up to 6,000 kWh/month and 0.30 above (Saudi Gazette 2025), repriced to 0.21 and 0.33 from January 2027 (Saudi Energy Company 2026). Compliance with the 2018 Saudi Building Code (Alayed et al. 2021) reduces energy use intensity by up to 20–25% in Climate Zone 2 (Olawale et al. 2025), but leaves plant operation unaddressed, and the rebound effect offsets part of the gains from appliance upgrades (Belaïd and Mikayilov 2024).

2.2 Control Approaches

The standard residential controller is a fixed-deadband thermostat with no view of tariff, occupancy or weather (Alotaibi et al. 2024). MPC optimises a cost over a rolling horizon, typically as a mixed-integer linear programme (Risbeck et al. 2017; Dinh et al. 2022), and reviews document substantial savings over rule-based control, largely under idealised conditions (Drgoňa et al. 2020). Its complexity grows with the number of binary variables and zones; for a single zone a dynamic-programming receding-horizon controller needs only a one-dimensional value grid (Section 9.4) and is used here as the causal model-based comparator.

RL learns a policy from interaction without an explicit model; most studies remain in simulation, with training cost, safety and generalisation the main barriers (Wang and Hong 2020). Reported savings range from 10 to 22% against rule-based or learned baselines (Azuatalam et al. 2020; Biemann et al. 2021; Du et al. 2021b,a) (Table 5); model-based, multi-agent and transfer variants extend the scope (Ding et al. 2025; Liu et al. 2024; Hou et al. 2024), and PPO matches SAC at lower cost in an empirical comparison (Manjavacas et al. 2024), which is the ground for adopting it here. Pure RL agents can violate comfort during exploration, and hybrid approaches impose hard rules that override the agent near the constraint boundaries (Xu et al. 2025).

2.3 Binary Actuation, Multi-Zone Coordination and Tariffs

Most RL-HVAC work targets continuous modulation (Wang et al. 2024), whereas Saudi residential split units are fixed-capacity and switched on and off, so a charge c_{sw} per switch-on stands for start-up wear; schedules that ignore it incur cycling and rebound costs the optimiser does not see (Faddel et al. 2023). For a single-variable Simple Linear Hybrid Automaton (SLHA) with per-mode running costs and per-switch costs, an optimal infinite-horizon schedule is computable in LogSpace while the finite-horizon problem is NP-complete, with a constant-factor approximation and an FPTAS (Mousa et al. 2016); the multi-variable extension is in NExpTime, with a PSpace ε-approximation and no hardness result (Mousa et al. 2017). In multi-zone buildings shared walls couple the zones: GA pre-scheduling with MADDPG reports 6.7% over a rule-based baseline (Xue et al. 2025), a Markov-game formulation is solved by multi-agent DRL (Yu et al. 2021), and in non-cooperative distributed MPC each zone minimises its own cost with no plant-wide optimality guarantee (Scattolini 2009). Most scheduling studies assume a constant price (El Makroum et al. 2023) or a smooth approximation (Risbeck et al. 2017); Proposition 1 below shows that a block tariff raises the price of the last kWh in a billing month but gives no incentive to rearrange consumption within it.

2.4 Research Gap

Table 1 summarises the positioning of the present work. The nearest neighbour to this study is the present author's own recent work in this journal (Faraj 2026). That paper addresses binary on/off scheduling for Saudi residential split units under the two-block tariff, with inter-zone thermal coupling, at two building scales, and with a cost-model sensitivity across step-wise, linear and exponential tariffs. It enforces comfort through a soft physics-informed reward and a top-k projection, neither of which excludes a violation, and it establishes no optimality certificate. It names as priority future work a quantification

of the remaining optimality gap against a day-ahead mixed-integer scheduler, a longer-horizon model-predictive baseline and modern constrained reinforcement learning, and separately, among its limitations, observes that formal constraint enforcement would provide stronger safety properties than its top-k projection heuristic.

The present study takes up both, though not by the route that paper anticipated. It supplies an optimality certificate, as a provable bound under the model that no schedule at least as cool as the baseline can beat (Section 8.1), and it compares receding-horizon controllers from 15 min to 48 h horizons under the same protocol as every other arm (Section 8.3). It enforces comfort by construction through a rule layer that overrides the learner's sampled action, a post-posed shield (Alshiekh et al. 2018), rather than through constrained policy optimisation. It does not quantify that paper's own optimality gap, which concerns a different instance, and it does not re-derive its cost-model or zone-count results, which stand. Its multi-zone evidence is narrower than that paper's: a Riyadh check, with independent per-room units, of whether the single-zone ordering survives in two- and four-room dwellings, without an optimality reference (Section 8.4).

Neither study addresses three points, which form the gap this paper fills: (i) hard comfort enforcement by construction, with the price of that guarantee measured; (ii) a clairvoyant optimality bracket against which a tuned rule, a learned policy and a receding-horizon controller are all read; and (iii) evaluation across three Saudi climates under one protocol, in which every arm is selected by the same rule and compared, week by week, with a baseline shifted to its own mean occupied temperature. Table 1 records two qualifications.

3 Problem Formulation

Closing those gaps requires a model in which comfort is a constraint rather than a reward term, and in which the price a schedule pays is the price the meter charges. This section builds that model in three parts: the thermal dynamics of a zone, the topologies over which they are defined, and the binary plant and tariff that price a schedule.

3.1 RC Thermal Model: Heat Loss, Gain, and Inter-Zone Transfer

Each zone $i \in \mathcal{N} = \{1, \dots, N_z\}$ is modelled as a first-order RC thermal node, the simplest member of the RC thermal-network family (Drgoňa et al. 2021; Wang et al. 2022; Shi et al. 2022) that is the standard control-oriented approach for residential buildings. Di Di Natale et al. (2023) demonstrated that physically consistent neural RC models scale to multi-zone settings while preserving interpretability; the present work retains the classical first-order RC for transparency and reproducibility. Let $C_{\text{th}}^{(i)}$ (kJ/K) be the effective thermal capacity of zone i, a single lumped value covering the room air together with its furnishings and the participating inner layers of the fabric (Remark 1), and $T_i(t)$ the temperature of that node; comfort is evaluated on it, so it stands for air and participating fabric together. The continuous-time energy balance for zone i is:

$$C_{\text{th}}^{(i)} \frac{\mathrm{d}T_i}{\mathrm{d}t} = \underbrace{-Q_{\text{hvac}}^{(i)} u_i(t)}_{\text{active cooling}} \underbrace{-\Lambda_i (T_i - T_{\text{out}})}_{\text{envelope}} + \underbrace{\sum_{j \in \mathcal{N}_i} \lambda_{ij} (T_j - T_i)}_{\text{inter-zone transfer}} + \underbrace{Q_{\text{sol}}^{(i)}(t)}_{\text{solar}}, \tag{1}$$

where $\Lambda_i = \lambda_{\text{ext}}^{(i)} + \lambda_{\text{win}}^{(i)}$ (W·K^{-1}) is the total envelope conductance, $u_i(t) \in \{0, 1\}$ is the binary HVAC state, and $\mathcal{N}_i$ is the set of adjacent zones.

Within one actuation interval the exogenous forcing (outdoor temperature, solar irradiance and the binary command) is held constant, so (1) is a linear system with constant inputs over the interval and its solution is available in closed form. Writing the system in deviation form $\dot{\mathbf{T}} = \mathbf{M}\left(\mathbf{T} - \mathbf{T}_\infty^{(t)}\right)$ with

$$\begin{aligned} \mathbf{M} &= -\tfrac{1}{C_{\text{th}}} [\operatorname{diag}(\Lambda_i + \Lambda_{\text{iz},i}) - [\lambda_{ij}]] , \\ \mathbf{T}_\infty^{(t)} &= -\mathbf{M}^{-1} \mathbf{f}^{(t)}, \end{aligned} \tag{2}$$

where $\mathbf{f}^{(t)}$ collects the forcing $C_{\text{th}}^{-1}\left(\Lambda_i T_{\text{out}}^{(t)} + Q_{\text{sol}}^{(i,t)} - Q_{\text{hvac}}^{(i)} u_i^{(t)}\right)$, the exact update over an interval of length Δt is

$$\mathbf{T}^{(t+1)} = \mathbf{T}_\infty^{(t)} + e^{\mathbf{M}\Delta t}\left(\mathbf{T}^{(t)} - \mathbf{T}_\infty^{(t)}\right), \tag{3}$$

Table 1: Positioning of the present study relative to selected prior work. ✓= addressed; ○ = partially addressed; – = not addressed.

Study	Binary on/off	Hard comfort	Step-wise tariff	Multi-zone	Saudi climate	Multi-city
(Azuatalam et al. 2020)	–	○	–	–	–	–
(Du et al. 2021b)	–	○	–	✓	–	–
(Biemann et al. 2021)	–	○	–	–	–	–
(Xu et al. 2025)	–	✓	–	–	–	–
(Xue et al. 2025)	–	○	–	✓	–	–
(Faraj 2026)	✓	○	✓	✓	✓	–
This work	✓	✓	○	✓	✓	✓

The step-wise tariff is formulated (Section 3.4), but the modelled dwelling never crosses a tier boundary under the schedule in force (Remark 3), hence ○. The multi-zone entry for this work covers one city and independent per-room units, without an optimality reference (Section 8.4). The ○ for hard comfort in (Faraj 2026) reflects enforcement by a soft physics penalty rather than by construction.

which for a single uncoupled zone reduces to the scalar form used throughout the single-zone analysis,

$$T^{(t+1)} = T_\infty^{(t)} + \left(T^{(t)} - T_\infty^{(t)}\right) e^{-\Delta t\, \Lambda_i / C_{\mathrm{th}}}, \qquad T_\infty^{(t)} = T_{\mathrm{out}}^{(t)} + \frac{Q_{\mathrm{sol}}^{(t)} - Q_{\mathrm{hvac}} u^{(t)}}{\Lambda_i}. \tag{4}$$

Equation (3) is used in preference to forward Euler, for three reasons. First, since $-\mathbf{M}$ is a weighted graph Laplacian plus the strictly positive diagonal $\Lambda_i/C_{\mathrm{th}}$, hence symmetric positive definite, every eigenvalue of $e^{\mathbf{M}\Delta t}$ lies in $(0, 1)$ for *every* $\Delta t > 0$: the update is unconditionally stable and no step-size stability condition arises. Forward Euler is stable only for $\Delta t < 2C_{\mathrm{th}}/\mu_{\max}$, with $\mu_{\max}$ the largest eigenvalue of $\mathrm{diag}(\Lambda_i) + [\Lambda_{\mathrm{iz},i}\delta_{ij} - \lambda_{ij}]$; at the air-only capacity of 77.2 kJ/K of the earlier model an hourly step met it for the single zone ($2C_{\mathrm{th}}/\Lambda_i = 1.8$ h) but not for the coupled topologies (Supplementary Section S1), so for the single zone the issue was accuracy and actuation, not stability. Second, the solution relaxes monotonically towards $\mathbf{T}_\infty^{(t)}$ within the interval, so for a single zone the trajectory's extremes are its endpoints and comfort can be checked exactly on the actuation grid; for coupled zones the check is on the grid only. Third, the step size affects only how often the controller may act, not the fidelity of the physics; the remaining restriction is the actuation condition (8), a property of the controller.

Inter-zone heat flux and effective inter-zone temperature.

The inter-zone coupling term in (3) is driven by a pairwise heat flux:

$$Q_{ij}^{(t)} = \lambda_{ij}\left(T_j^{(t)} - T_i^{(t)}\right), \tag{5}$$

positive when zone j is warmer. With net load $Q_{\mathrm{iz},i}^{(t)} = \sum_j Q_{ij}^{(t)}$, total inter-zone conductance $\Lambda_{\mathrm{iz},i} = \sum_j \lambda_{ij}$ and the *effective inter-zone temperature*

$$T_{\mathrm{iz},i}^{(t)} = \frac{\sum_{j \in \mathcal{N}_i} \lambda_{ij}\, T_j^{(t)}}{\Lambda_{\mathrm{iz},i}}, \tag{6}$$

the net inter-zone load simplifies to $Q_{\mathrm{iz},i}^{(t)} = \Lambda_{\mathrm{iz},i}\left(T_{\mathrm{iz},i}^{(t)} - T_i^{(t)}\right)$, exactly mirroring the envelope term $\Lambda_i(T_i^{(t)} - T_{\mathrm{out}}^{(t)})$. In a multi-zone instance this variable is the natural state input for anticipating inter-zone heat cascades; the multi-zone receding-horizon controller of Section 8.4 uses its neighbours' planned temperatures in this form, and no learned policy observes it here.

The RC model serves both as the simulation environment and as the prediction model of the

receding-horizon controller (Section 5.5); its state is available to the PPO agent through the state vector (Section 5.2), but the trained policy does not use it to plan. Table 2 states the assembly transmittances and areas from which the conductances are formed. The two transmittances are at the SBC 2018 prescriptive maximum (Alayed et al. 2021) of $U_{\text{wall}} = 0.45$ and $U_{\text{win}} = 2.8\,\text{W}\,\text{m}^{-2}\text{K}^{-1}$, giving a fabric conductance of 11.90 W/K per exposed face. Section 6.2 relates the conductances to the code. The code constrains assemblies and not a zone's total conductance, so the compliance claim is made at assembly level only. Infiltration is carried explicitly at 0.6 air changes per hour and evaluated per city from station pressure, contributing 10.77–12.72 W/K (slightly more than one face of fabric). Compliance is mandatory and enforced through municipal permitting and, since October 2024, the activation of the electricity connection (Argaam 2024).

The occupied comfort band is $[T_{\min}, T_{\max}]$=[22, 24] °C; its provenance is set out in Section 6.2.

Remark 1 (Thermal capacity, the actuation condition, and stability) The RC model (3) follows the control-oriented identification framework of Bacher and Madsen (2011) and captures sensible heat only. The capacity is $C_{\text{th}} = 4{,}800\,\text{kJ/K}$, an areal heat capacity of $50\,\text{kJ}\,\text{m}^{-2}\text{K}^{-1}$ over $96\,\text{m}^2$ of participating surface, within the ISO 13786 range for a furnished heavy-construction room; with $\Lambda_i = 23.75\,\text{W/K}$ the time constant is $C_{\text{th}}/\Lambda_i = 56\,\text{h}$. ISO 13786's areal capacity is a 24 h quantity, and the capacity effective at the period of THERM's cycle (about 6 h) is lower, so the reference value probably lies above the effective one, towards the capacities at which the tuned thermostat gains least; the results are therefore repeated at 2,400, 7,200 and 9,600 kJ/K (Section 8.3).

The actuation interval is constrained twice over. Holding the zone inside the band across one interval requires the plant's delivery over that interval to be no larger than the band:

$$\frac{Q_{\text{hvac}}\,\Delta t}{C_{\text{th}}} \;\le\; T_{\max} - T_{\min}. \tag{7}$$

This follows from the energy balance and is *invariant to the integration scheme*: subdividing the interval delivers the same energy and the same excursion, so a finer integrator cannot remedy a violation. But (7) is only necessary. A controller switching on a hysteresis deadband must complete each excursion inside *its own* deadband δ_{db}, and the deadband together with those excursions must in turn fit inside the comfort band, since the zone keeps moving for one interval after each switch is commanded. With ϵ_{hot} the warming excursion of one interval at peak load (19), both requirements give

$$\frac{Q_{\text{hvac}}\,\Delta t}{C_{\text{th}}} \;\le\; \delta_{\text{db}} \;\le\; (T_{\max}-T_{\min}) - \frac{Q_{\text{hvac}}\,\Delta t}{C_{\text{th}}} - \epsilon_{\text{hot}}, \tag{8}$$

and for a deadband centred on the band the upper limit tightens to $2\min\!\big(\tfrac{1}{2}(T_{\max}-T_{\min}) - Q_{\text{hvac}}\Delta t/C_{\text{th}},\ \tfrac{1}{2}(T_{\max}-T_{\min}) - \epsilon_{\text{hot}}\big)$. At the parameters of Table 2 the admissible interval is $0.375 \le \delta_{\text{db}} \le 1.25\,°\text{C}$, against the 2 °C that the band width alone would allow. The distinction matters: a 1.5 °C deadband satisfies the band-width bound yet leaks 0.132 occupied degree-hours, so the looser form cannot adjudicate the admissibility sweeps it is used for. The THERM baseline of Section 5.5 switches at $T_{\max} - 0.5$ and $T_{\min} + 0.5$, so $\delta_{\text{db}} = 1\,°\text{C}$ on the [22, 24] °C band: inside the interval admissible at a 15-min step and near its upper end (Section 9.2).

This determines the interval used throughout. At $\Delta t = 15\,\text{min}$ the left-hand side is 0.375 °C against $\delta_{\text{db}} = 1\,°\text{C}$, a factor of 2.7. At $\Delta t = 1\,\text{h}$ it is 1.50 °C and the condition fails: THERM then accumulates 4.47 occupied degree-hours in the first held-out Riyadh week, so the two arms are no longer held to the same constraint. At 30 min, and equally at 2,400 kJ/K and 15 min, the upper limit of (8) is 0.99 °C and fails narrowly. With the 690 W peak load of (19), an hourly interval would require $C_{\text{th}} \ge (Q_{\text{hvac}} + 690\,\text{W})\Delta t/(T_{\max} - T_{\min} - \delta_{\text{db}}) \approx 9{,}700\,\text{kJ/K}$ by (8), and 14,400 kJ/K for a deadband centred on the band as THERM's is, against the 3,600 kJ/K that (7) alone would suggest, so it is inadmissible here and $\Delta t = 15\,\text{min}$ is adopted. The admissibility sweeps over interval and capacity and the validation literature for the RC form are in Supplementary Section S1 and Table S1; the dependence of the saving on the interval is reported in Section 9.3.

Remark 2 (What scheduling can and cannot change) One structural property of the plant bounds every claim made below. Summing (3) over zones, each unordered pair $\{i, j\}$ contributes $\lambda_{ij}(T_j - T_i) + \lambda_{ij}(T_i - T_j) = 0$, so the inter-zone terms cancel exactly: conduction between zones redistributes heat but adds none. Averaging over a periodic steady state, in which the mean of $\mathrm{d}T_i/\mathrm{d}t$ vanishes, leaves

$$Q_{\text{hvac}}\,\langle u_i\rangle = \Lambda_i\big(\langle T_{\text{out}}\rangle - \langle T_i\rangle\big) + \langle Q_{\text{sol}}^{(i)}\rangle, \tag{9}$$

so the duty cycle, and hence the daily electrical energy $E_{\text{day}} = 24\,P_{\text{hvac}}\langle u_i\rangle$, is determined by the mean zone temperature alone, given the weather. With $\Lambda_i = 23.75\,\text{W/K}$ for the single-face Riyadh zone and a fixed COP, $\partial E_{\text{day}}/\partial\langle T_i\rangle = -0.191\,\text{kWh/day/K}$.

This constrains the interpretation of the results in three ways. First, a controller confined to the 2 °C occupied band can move $\langle T_i \rangle$ by at most 1 K from the band centre, bounding the energy it can save relative to a thermostat there at about 0.19 kWh/day (0.38 across the full band). Second, with a constant COP and an occupied band that never changes, pre-cooling *lowers* the mean temperature and therefore *raises* consumption; an occupancy setback, which changes the band in time, or a temperature-dependent COP would alter this. Third, a cost comparison between two controllers at different mean temperatures is largely a restatement of (9), which is why every arm in Section 8 is compared with a baseline shifted to its own mean occupied temperature (Section 7).

At equal mean temperature every arm therefore uses the same terminal-priced energy, up to the finite-window and non-periodic residual of the identity, and the saving equals c_{sw} times the on-commands avoided. With a constant COP and no start-up energy this is an identity of the model, not a finding: the cost saving reported here is a charge on avoided on-commands, assumed to stand for deferred wear, neither an energy saving nor a reduction in the electricity bill, and the energy and on-command columns are reported beside it.

3.2 Zone Configurations

Three topologies are defined. The headline study evaluates the single zone; Section 8.4 evaluates two- and four-room instances in Riyadh; the topologies are defined in the repository code (`mz_r3.py`, the Data availability statement). The *single zone* (1×1) has one windowed exterior wall and no inter-zone coupling. The *linear array* ($1 \times N$, $N \in \{2, 4, 6\}$) places N rooms in a row, each with one windowed exterior wall (azimuths in Table S-MZ); interior rooms couple to two neighbours. The *grid* ($N \times N$, $N \in \{2, 3, 4\}$) produces up to $N_z = 16$ zones; corner zones have two windowed exterior walls, edge zones one.

3.3 Binary HVAC Model and Interval Cost

Each zone has an independent split-unit air conditioner with two modes: *on* ($u_i = 1$) or *off/idle* ($u_i = 0$, at zero operational cost). The interface is supervisory: $u_i = 1$ commands the unit to deliver its full sensible capacity for the 15-min step, as a smart plug or an infrared command at a low set-point would, and the model lumps air and mass in one node, so the unit's own return-air cycling is not represented. A charge c_{sw} (SAR) is incurred on each $0 \to 1$ transition, an *on-command*, standing for start-up wear. Whether each on-command corresponds to a compressor start depends on the equipment's own thermostat and is not established here, so the events counted and priced throughout are on-commands. Switching off carries no charge.

Definition 1 (Interval cost) The cost of zone i at interval t under model m is:

$$c_i^{(t)} = c_{\mathrm{sw}} \cdot \mathbf{1}\big[u_i^{(t)} > u_i^{(t-1)}\big] + f^m\big(u_i^{(t)}, E_{\mathrm{cum}}^{(t)}\big) \cdot \Delta t, \tag{10}$$

where $E_{\mathrm{cum}}^{(t)} = \sum_{\tau \le t} \sum_i P_{\mathrm{hvac}} u_i^{(\tau)} \Delta t$ is cumulative billing-month energy (kWh) up to interval t.

3.4 Electricity Cost Models

Model L (linear) charges a constant rate, the standard assumption in scheduling (El Makroum et al. 2023): $f^{\mathrm{L}} = c_r^{\mathrm{L}} P_{\mathrm{hvac}} u_i$. **Model E** (exponential), $f^{\mathrm{E}} = c_r^{\mathrm{E}} e^{\beta E_{\mathrm{cum}}} P_{\mathrm{hvac}} u_i$ with $\beta = 5 \times 10^{-4}\,\mathrm{kWh}^{-1}$, is a smooth approximation to step-wise pricing over a comparable price range. Neither is used in any result.

Model S (Step-wise Cost Model). Model S is the step-wise (piecewise-constant) cost model. It is instantiated with the Saudi four-tier residential schedule in force from January 2016 to December 2017 (Nahiduzzaman et al. 2023), in which the marginal price rises across four consumption blocks. This is a historical schedule, no longer the tariff in force, and it is used here as a *stress-test* instantiation of the step-wise cost class, not a representation of current Saudi prices; Remark 3 treats the tariff now in force separately and quantifies what it implies for the modelled dwelling. The billed quantity E_{cum} is the whole-house monthly cumulative consumption (the non-HVAC background load of Section 7 plus HVAC); VAT and the fixed meter charge are excluded as they do not affect scheduling:

$$p(E_{\mathrm{cum}})\ [\mathrm{SAR/kWh}] = \begin{cases} 0.05, & E_{\mathrm{cum}} \le 2{,}000, \\ 0.10, & 2{,}000 < E_{\mathrm{cum}} \le 4{,}000, \\ 0.20, & 4{,}000 < E_{\mathrm{cum}} \le 6{,}000, \\ 0.30, & E_{\mathrm{cum}} > 6{,}000. \end{cases} \tag{11}$$

Table 2: Zone parameters for each $4 \times 4 \times 4$ m room.

Symbol	Description	Value	Unit
C_{th}	Thermal capacity (furnished; ISO 13786, Remark 1)	4,800	kJ/K
U_{wall}	Opaque-wall transmittance (SBC 2018 prescriptive max)	0.45	$\text{W}\cdot\text{m}^{-2}\text{K}^{-1}$
A_{wall}	Opaque area per exterior wall	14.0	m^2
U_{win}	Window transmittance (SBC 2018 prescriptive max)	2.8	$\text{W}\cdot\text{m}^{-2}\text{K}^{-1}$
A_{gl}	Glazing area per exterior wall	2.0	m^2
$\lambda_{\text{ext}} + \lambda_{\text{win}}$	Fabric conductance *per exposed face*	11.90	$\text{W}\cdot\text{K}^{-1}$
Λ_{inf}	Infiltration conductance at 0.6 ACH (per city)	10.77–12.72	$\text{W}\cdot\text{K}^{-1}$
SHGC	Solar heat-gain coefficient	0.25	—
E_{base}	Non-HVAC monthly household load	800	kWh/month
λ_{ij}	Inter-zone shared-wall conductance	12.5	$\text{W}\cdot\text{K}^{-1}$
Q_{hvac}	HVAC sensible extraction (assumed)	2.0	kW
COP	Coefficient of performance (assumed, constant)	2.985	—
P_{hvac}	HVAC electrical draw	0.670	kW
$T_{\min}/T_{\max}$	Occupied comfort band	22/24	°C
δ_{db}	THERM deadband (baseline, Section 5.5)	1.0	°C
$\epsilon_{\text{hot}}/\epsilon_{\text{cold}}$	Rule guard bands, from (19)	0.20/0.40	°C
c_{sw}	Charge per on-command (computational reference)	0.15	SAR
Δt	Actuation interval	15	min

Envelope conductance is affine in the number of exposed faces n, $\Lambda_i(n) = 11.90\,n + \Lambda_{\text{inf}}$, because infiltration is a property of the zone volume. Λ_{inf} depends on station pressure: 11.85 W/K at Riyadh (94.3 kPa), 12.72 at Jeddah (101.2 kPa) and 10.77 at Taif (85.7 kPa, 1,478 m), so the single-face totals are $\Lambda_i = 23.75$, 24.62 and 22.67 W/K.

Proposition 1 (The bill does not depend on timing) *Let the marginal price be any locally integrable function $p(\cdot)$ of cumulative billing-month consumption alone, and let E_{cum} be non-decreasing and absolutely continuous in t. Then the energy bill over the billing month is $B = \int_0^{E_{\text{total}}} p(e)\,\mathrm{d}e$, which depends only on E_{total} and not on the order in which it was consumed.*

Proof Charging the rate in force as consumption accumulates gives $B = \int_0^T p\big(E_{\text{cum}}(t)\big)\dot{E}_{\text{cum}}(t)\,\mathrm{d}t$. Absolute continuity makes $e = E_{\text{cum}}(t)$ a valid change of variable and non-decrease makes it orientation-preserving, so $B = \int_0^{E_{\text{total}}} p(e)\,\mathrm{d}e$. No property of p beyond local integrability is used; in particular monotonicity is not required. □

On the actuation grid the realised bill is path-dependent only when a tier boundary falls inside one interval's increment, at the order of one interval's consumption (Supplementary Section S1; no boundary is approached here, Remark 3).

Three consequences follow. First, Models L, E and S are all of this form (constant, exponential in E_{cum}, and a step function of E_{cum}), so none of them can be distinguished by any experiment that rearranges consumption in time. What separates them is the exchange rate $c_{\text{sw}}/(pP_{\text{hvac}}\Delta t)$ between an avoided switch and an hour of running: a constant under Model L, and under Models E and S a quantity that varies with E_{cum} through the billing month, by a factor of only 1.04 under Model E between the background load and this account's monthly consumption. The three therefore induce a one-parameter family of behaviours whose range Section 9.2 sweeps, and a separate cost-model ablation would resolve nothing those sweeps do not. Second, pre-cooling is dominated under *every* such tariff, not only under a constant price: it changes when energy is drawn and, by identity (9), raises how much, so the timing earns nothing and the extra energy costs. That is why the rule layer contains no pre-cooling rule (Section 5.3). Third, the proposition and its second consequence fail under different conditions. The *proposition* fails for any price that is not a function of cumulative consumption alone (a demand or capacity charge on coincident power, time-of-use and real-time pricing) and, trivially, across a billing boundary, since the tiers reset monthly and shifting load between months is a genuine arbitrage. The second *consequence*, that pre-cooling is dominated, fails under two further conditions that leave the proposition intact: a minimum on- or off-time constraint, which makes the

cost of a given duty depend on its arrangement, and a temperature-dependent coefficient of performance, which is a property of the plant, not of the tariff, and which would make cooling genuinely cheaper at night. Section 10.4 records that none of these is evaluated here, and that the conclusions of this paper do not extend to them.

The four-tier schedule is retained as Model S because it has three tier boundaries within the monthly consumption of a larger dwelling and is therefore the most demanding instance of the step-wise class. Two properties matter. First, the marginal price is a non-decreasing step function of cumulative consumption, so the bill is convex and representable in a linear programme; the difficulty lies in the binary actions and the switching term, not the tariff. Second, under Model S the per-step cost rate depends on the accumulated state E_{cum}, which takes the problem outside the SLHA class of Section 4.2 and requires E_{cum} in the agent state. The choice does *not* establish performance under present-day Saudi prices.

Remark 3 (The tariff currently in force, and what it implies here) The residential tariff in force from January 2018 to December 2026 is a two-block schedule: 0.18 SAR/kWh up to 6,000 kWh/month and 0.30 SAR/kWh above (Saudi Gazette 2025). From 1 January 2027 the blocks are repriced to 0.21 and 0.33 SAR/kWh and a fixed monthly charge of 70–100 SAR, graded by breaker capacity, replaces the former meter reading, maintenance and billing fee (Saudi Energy Company 2026). Both have a single boundary at 6,000 kWh/month. The fixed charge is schedule-independent and does not alter the optimal policy, but it enlarges the bill and so reduces every percentage saving quoted here. Both belong to the step-wise class of (11).

The modelled account consumes ≈884 kWh per billing month (Section 7), far below the boundary, so it stays in the first block throughout and faces a *constant* marginal price. The scheduling problem therefore belongs to the same cost class as Model L, though not the same instance: Model L is specified at 0.12 SAR/kWh (no result uses it) against a first-block rate of 0.18 (0.21 from 2027). Because c_{sw} is a fixed charge that does not scale with the energy price, the ratio $c_{\mathrm{sw}}/(p\,P_{\mathrm{hvac}}\,\Delta t)$, which determines the optimal switching policy, falls by a third between those rates. Every arm is therefore reported at the rate in force, and TUNED and the bound are repeated at the announced 2027 rate (Section 9.2; the receding-horizon controller is not). There the bound on what a schedule no warmer than the baseline can save in Riyadh is 29.6, 20.0 and 18.6% across the three schedules, tracking the on-command share of the baseline's cost, 81.5, 55.0 and 51.2%.

The saving therefore arises from longer, less frequent on-phases, that is, from the switching condition (O3 of Section 4.4), not from the energy condition (O2).

Running example. A single-zone instance (1×1, Riyadh, July, $T_{\mathrm{out}} = 40\,^{\circ}\mathrm{C}$) at the parameters of Table 2 is used throughout to fix intuition about the cycle structure of Observation 2; Figure 1 draws it from the closed-form scalar solution. ◇

With the thermal and cost models established, the following section formalises the constrained optimisation and its optimality conditions.

4 Optimal Schedule: Constrained Formulation and Optimality Conditions

This section casts the problem over an infinite horizon (Section 4.1), relates its single-zone core to simple linear hybrid automata (Section 4.2), states the monthly constrained problem (Section 4.3) and derives the optimality conditions and a corollary on the rule layer (Section 4.4).

4.1 Infinite Time Horizon and Limit-Average Cost

A residential HVAC system operates indefinitely: months repeat and no planning horizon ends the building's operation, so the relevant cost measure is the long-run average cost per interval. Mousa et al. (2016), the problem is therefore cast as an *infinite time horizon* optimisation whose performance criterion is the *limit-average cost*.

Definition 2 (Infinite schedule and limit-average cost) An *infinite schedule* is the binary matrix $\sigma = (u_i^{(t)})_{N_z\times\infty}$, $u_i^{(t)} \in \{0,1\}$. Its limit-average cost under model m is:

$$\bar{J}_{\infty}^{m}(\sigma) = \limsup_{H\to\infty} \frac{1}{H} \sum_{t=0}^{H-1} \sum_{i\in\mathcal{N}} c_i^{(t)}, \qquad (12)$$

where $c_i^{(t)}$ is the interval cost of Definition 1. A schedule is *safe* if, for all i and all $t \geq 0$, $T_i^{(t+1)} \leq$

$T_{\max}$ on every occupied step and $u_i^{(t)} = 1$ implies $T_i^{(t+1)} \geq T_{\min}$: the ceiling is pointwise and a cooling step may never end below $T_{\min}$, while a free-floating step ($u_i^{(t)} = 0$) has no lower limit.

Every number reported is computed on a finite horizon: each schedule is scored over 168 h windows of $N = 672$ steps, the receding-horizon controller re-plans every step over a look-ahead of 0.25–48 h, and the clairvoyant schedule and the bound treat each window as a whole, with the terminal value V_N (Section 5.5) standing in for the continuation so that the end of a window is not exploited.

4.2 Simple Linear Hybrid Automata: Formal Connection

The SLHA connection (Mousa et al. 2016) provides a formal language for *complete cycles*, establishes that an optimal infinite-horizon schedule for the single-zone constant-tariff problem is computable in LogSpace, and identifies which features of the full problem take it outside that tractable class.

Definition 3 (Single-zone SLHA) The single-zone binary HVAC system under Model L is the SLHA $\mathcal{A} = (Q, X, \mathrm{Inv}, f, \Delta, \mathcal{C})$ with modes $Q = \{q_0 \text{ (idle)}, q_1 \text{ (cool)}\}$, variable $X = \{T\}$, flows $\dot{T}{=}a_0 \triangleq (\Lambda(T_{\text{out}} - T) + Q_{\text{sol}})/C_{\text{th}}$ in q_0 and $\dot{T}{=}a_1 \triangleq a_0 - Q_{\text{hvac}}/C_{\text{th}}$ in q_1 ($a_1 < 0 < a_0$ in Saudi summer), invariants $T \leq T_{\max}$ in q_0 and $T \geq T_{\min}$ in q_1, transitions $(q_0 \to q_1)$ on $T \geq T_{\max} - \epsilon_{\text{hot}}$ with cost c_{sw} and $(q_1 \to q_0)$ on $T \leq T_{\min} + \epsilon_{\text{cold}}$ with cost 0, and continuous cost rate $c_r P_{\text{hvac}}$ in q_1 only.

A state diagram of the automaton is in the repository (the Data availability statement).

The rates of Definition 3 depend on T through $\Lambda(T_{\text{out}} - T)$, so reading them as constants is an abstraction; at $\tau = 56$ h its error in the band-crossing time is 0.1–0.2% (Supplementary Section S2), and every numerical result is computed on the exact dynamics (3).

Multi-zone coupling, time-varying disturbances and Model S's consumption-dependent cost rate each take the problem outside the SLHA class, since the LogSpace result of (Mousa et al. 2016) requires a single continuous variable, a time-invariant automaton and a constant per-mode cost rate. No polynomial-time algorithm is known for the extended problem, and the finite-horizon single-variable case is already NP-complete (Mousa et al. 2016); no hardness result is claimed for the multi-zone problem, which remains open.

Definition 4 (Leap and complete cycle) A *leap* is a maximal contiguous time interval during which the automaton stays in one mode. An *on-leap* ℓ^+ (in mode q_1) of duration τ^+ cools the zone from $T_{\max}$ to $T_{\min}$ at rate $|a_1|$: $\tau^+ = (T_{\max} - T_{\min})/|a_1|$. An *off-leap* ℓ^- (in mode q_0) of duration τ^- returns temperature from $T_{\min}$ to $T_{\max}$ at rate a_0: $\tau^- = (T_{\max} - T_{\min})/a_0$. The discrete switch-on cost c_{sw} is incurred at the start of each on-leap only. A *complete cycle* is one on-leap followed by one off-leap. Its total cost and duration are:

$$\mathcal{C}(\ell^+, \ell^-) = c_{\text{sw}} + c_r P_{\text{hvac}}\, \tau^+, \quad \tau_{\text{cyc}} = \tau^+ + \tau^-. \tag{13}$$

The resulting *limit-average cost per unit time* of repeating complete cycles indefinitely is:

$$\bar{c}^{\,*} = \frac{c_{\text{sw}} + c_r P_{\text{hvac}}\, \tau^+}{\tau^+ + \tau^-}. \tag{14}$$

Observation 1 (Optimal infinite schedule for SLHA, Mousa et al. (2016)) The optimal safe infinite schedule for the single-zone SLHA of Definition 3 is to repeat the cheapest complete cycle (ℓ^+, ℓ^-) indefinitely. Its limit-average cost equals $\bar{c}^{\,*}$ of equation (14), which is achievable in deterministic LogSpace (Theorem 1 of (Mousa et al. 2016)).

The SLHA analysis is not used as a solution method for the full problem but as a diagnostic tool that identifies the structural properties on which the controller design draws. It yields three design inputs: (i) optimal schedules have complete-cycle structure (Observation 2), motivating rules R1–R2; (ii) the ratio τ^+/τ^- governs cycle structure and motivates the one-hour outdoor forecast in the agent state; (iii) the absence of any known polynomial-time algorithm for the multi-zone problem (Mousa et al. 2017) motivates approximate controllers, learned or receding-horizon. No controller here solves the monthly problem (15) itself: it states the objective, in line with the tariff's monthly reset, that the controllers approximate on a receding basis and that is evaluated on 168 h windows.

The rule layer is weaker than the structural result that motivates insight (i). Observation 2 characterises *optimal* single-zone schedules as sequences of complete cycles. Rules R1 and R2 of Definition 7, by contrast, are boundary-triggered overrides: they fix the action only on $T_i \geq T_{\max} - \epsilon_{\text{hot}}$ and on $T_i \leq T_{\min} + \epsilon_{\text{cold}}$, and leave it unconstrained on the interior $(T_{\min} + \epsilon_{\text{cold}}, T_{\max} - \epsilon_{\text{hot}})$, which for the guard bands of (19) is 1.40 K of the 2 K band. They impose no minimum dwell and never require a cycle, once begun, to complete, so a schedule switching every step inside the interior satisfies them exactly with no complete-cycle structure at all. The rule layer thus enforces O1 by construction (for the nominal plant parameters from which (19) derives its margins) and is *consistent with* the structure of Observation 2 without *implementing* it; complete cycles are left to the learned policy's switching term, with no guarantee. A rule set that did implement it (a minimum-dwell constraint) is identified as required work in Section 10.4.

4.3 Constrained Optimisation Problem

Definition 5 (Schedule over one billing month) A *monthly schedule* over horizon $\mathcal{H} = \{0, \ldots, H-1\}$ (with $H = 2{,}880$ intervals of 15 min for a 30-day month) is the binary matrix $\sigma = (u_i^{(t)})_{N_z \times H}$, $u_i^{(t)} \in \{0, 1\}$.

Definition 6 (Constrained optimisation problem) The monthly scheduling problem is:

$$
\begin{aligned}
\underset{\sigma \in \{0,1\}^{N_z \times H}}{\text{minimise}} \quad & \bar{J}^m(\sigma) = \frac{1}{H} \sum_{t \in \mathcal{H}} \sum_{i \in \mathcal{N}} c_i^{(t)} \\
\text{subject to} \quad & T_i^{(t+1)} \leq T_{\max} \quad \forall\, i,\, t \in \mathcal{H}_{\text{occ}} \\
& \text{(ceiling, } \textbf{hard}\text{)} \\
& u_i^{(t)} = 1 \Rightarrow T_i^{(t+1)} \geq T_{\min} \quad \forall\, i,\, t \in \mathcal{H} \\
& \text{(cold rule, } \textbf{hard}\text{)} \\
& T_i^{(t+1)} = f\Big(T_i^{(t)}, u_i^{(t)}, \{T_j^{(t)}\}_{j \in \mathcal{N}_i}, T_{\text{out}}^{(t)}\Big) \\
& \text{(RC dynamics, eq. (3))}
\end{aligned}
\tag{15}
$$

Here $\mathcal{H}_{\text{occ}} \subseteq \mathcal{H}$ is the set of occupied steps. A schedule satisfying both comfort constraints is *safe* (Definition 2). Both are *hard*: no cost saving justifies an occupied step ending above $T_{\max}$ or a cooling step ending below $T_{\min}$. A free-floating step has no lower limit, because a cooling-only plant cannot stop a zone drifting below $T_{\min}$ when the outdoor air is cool; the protocol (Section 7) reports such free-float undertemperature and gives it no credit in the clipped mean, rather than counting it as a violation of the schedule. On away steps, used only by the setback arms, the ceiling is the empty-house limit of Section 8.2. The RC dynamics embed four heat-transfer mechanisms (active cooling, envelope loss/gain, solar heat gain, and bidirectional inter-zone transfer), making this a multi-dimensional, non-linear, constrained binary programme. The inter-zone term $\sum_{j \in \mathcal{N}_i} \lambda_{ij}(T_j - T_i)$ is embedded directly in the RC constraint (3): heat conducted from a warmer neighbour raises $T_i^{(t+1)}$, potentially triggering the ceiling constraint in zone i even when $u_i^{(t)} = 0$. An optimal schedule must therefore coordinate all zones simultaneously. The multi-zone study of Section 8.4 has no optimality reference, so it does not measure what that coordination is worth.

4.4 Optimality Conditions

What makes a schedule optimal?

A safe schedule σ^* is *optimal* if it satisfies O1 and resolves the O2–O3 trade-off at least cost, where the three conditions are: **(O1) Comfort feasibility**: $T_i^{(t)} \leq T_{\max}$ for all i and every occupied t, and no cooling step ends below $T_{\min}$; both are hard, and free-float drift below the band is measured (Definition 6). **(O2) Minimum energy at the prevailing marginal price**: the schedule draws as little energy as the comfort constraint permits. Under every cost model used here this is the whole of the economic content of running time, because Proposition 1 shows the bill depends on total consumption and not on its arrangement; there is no "cheapest window" to concentrate into. **(O3) Minimum switching frequency**: on-phases are as long and infrequent as the thermal state permits, amortising the discrete cost c_{sw} per activation.

O1 enforces the hard constraint; O2 and O3 are the two economic terms of $\bar{J}^m$ in equation (15), and they can trade off. At a fixed mean temperature the energy term is fixed by identity (9), so what separates schedules compared at equal mean is O3 (Section 8.2). In the infinite-horizon interpretation (12), an optimal schedule repeats a cost-minimising monthly pattern indefinitely. For the single-zone SLHA subproblem this pattern is the cheapest complete cycle (Observation 1). No

claim is made that the learned schedules attain it (Section 4.2).

For the running example ($T_{\text{out}} = 40\,^{\circ}\text{C}$, flat price at the in-force 0.18 SAR/kWh), the complete cycle of Figure 1 gives $\bar{c}^{*} \approx 0.042$ SAR/h; switching every 3 h instead, at the same duty ratio and so violating O3, costs 0.074 SAR/h, which is 75% more. That difference is O3's monetary value.

Remark 4 (Switching cost and schedule granularity) Condition O3 implies that a higher c_{sw} favours fewer, longer on-phases (Observation 2), and therefore that the saving available to any scheduler grows with it. Section 9.2 measures that dependence and Section 10.3 draws the consequence: the reported figure is a property of the tariff as much as of the controller. The guard thresholds do not need to move with c_{sw}, because (19) derives them from the actuation interval and the load, neither of which depends on a price.

Observation 2 (Complete-cycle structure) Within the SLHA abstraction of Section 4.2 (single zone, $c_{\text{sw}} > 0$, constant marginal price), an optimal schedule consists of *complete cycles*: an on-phase during which T_1 cools from $T_{\max}$ to $T_{\min}$, followed by an idle phase during which T_1 drifts back to $T_{\max}$. This is Theorem 2 of (Mousa et al. 2016) applied to the abstraction: a schedule that switches on before $T_{\max}$ or off before $T_{\min}$ can be merged into a longer on-phase at the same continuous cost and one fewer switch. The corollary below does not rely on it.

Figure 1 draws this structure at the parameters of this study. Observation 2 has a consequence for a learner placed inside the rule layer, under hypotheses that the evaluated problem does not meet; the corollary states it and Remark 5 measures what the gap between the two is worth.

Corollary 1 (Optimality of the rule-induced hysteresis under constant disturbance) *Consider the single-zone problem with $c_{\text{sw}} > 0$ under a constant marginal price, with the hard rules of Definition 7 in force,* with T_{out} and Q_{sol} constant*, so that the drift rates a_0, a_1 of Definition 4 are time-invariant, and compare policies* at the mean occupied temperature of the maximal-width cycle*, the midpoint of $\mathcal{F}$ below. Let $\mathcal{F} = (T_{\min}+\epsilon_{\text{cold}},\, T_{\max}-\epsilon_{\text{hot}})$ be the interval on which the agent's action is not overridden: the rules fix the action everywhere else, so this is the whole class over which a learned policy can act. Then the rule-induced hysteresis (continue the previous action throughout $\mathcal{F}$) attains cost no greater than that of any policy that differs from it on $\mathcal{F}$.*

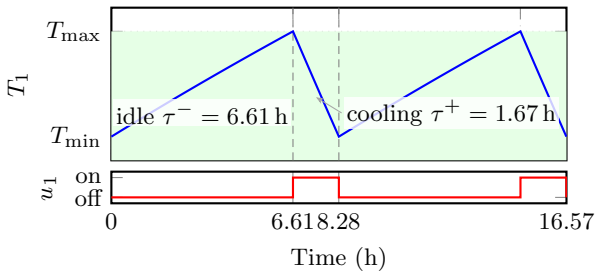


Fig. 1: The complete-cycle structure of Observation 2, drawn from the closed-form scalar solution at $T_{\text{out}} = 40\,^{\circ}\text{C}$, $\Lambda_i = 23.75$ W/K and $C_{\text{th}} = 4800$ kJ/K. The optimal schedule touches each band edge exactly once per cycle: the plant runs only while cooling $T_{\max} \to T_{\min}$ and is idle while the zone drifts back. The dwell times follow in closed form from the scalar update, $\tau^- = \tau \ln[(T_{\text{out}}-T_{\min})/(T_{\text{out}}-T_{\max})] = 6.61$ h and $\tau^+ = \tau \ln[(T_{\max}-T_{\infty}^{+})/(T_{\min}-T_{\infty}^{+})] = 1.67$ h with $\tau = C_{\text{th}}/\Lambda_i = 56.1$ h, giving a cycle of 8.28 h, or 2.90 switch-ons per day at this constant $40\,^{\circ}\text{C}$ load; on the held-out Riyadh weeks the clairvoyant schedule issues 2.83 on-commands per day (THERM's 4.10 less the 31.1% of Table 3). This geometry is what Corollary 1 rests on.

Proof Write $T_{\text{on}}^{g} = T_{\max} - \epsilon_{\text{hot}}$ and $T_{\text{off}}^{g} = T_{\min} + \epsilon_{\text{cold}}$ for the guard thresholds, and $W = T_{\text{on}}^{g} - T_{\text{off}}^{g} = |\mathcal{F}|$ for the width they enclose.

Step 1: phases and widths. Any policy is a sequence of on-phases and off-phases. On-phase k starts at U_k and ends at L_k; off-phase k starts at L_k and ends at U_{k+1}. The rules force the action outside $\mathcal{F}$, so every switch point lies in $[T_{\text{off}}^{g}, T_{\text{on}}^{g}]$ and the widths $w_k = U_k - L_k$ and $v_k = U_{k+1} - L_k$ obey $w_k, v_k \le W$. No merging argument is needed.

Step 2: the energy term agrees across all competitors. With a_0, a_1 constant, on-phase k lasts $w_k/|a_1|$ and off-phase k lasts v_k/a_0. Because the temperature stays in a bounded interval, $|\sum_{k\le n} w_k - \sum_{k\le n} v_k| \le W$ for every n, so every policy has the long-run duty ratio $a_0/(a_0 + |a_1|)$ (identity (9) in its time-invariant form) and the same energy cost per unit time at a constant marginal price. The comparison reduces to the switching term.

Step 3: switch count is minimised by maximal width. Over a horizon H containing $n(H)$ starts, $H = \sum_{k\le n(H)} \big(w_k/|a_1| + v_k/a_0\big) + O(W/a_0)$, so

$$\lim_{H\to\infty} \frac{n(H)}{H} \;=\; \Big[(1/a_0 + 1/|a_1|)\;\overline{w}\Big]^{-1}, \qquad \overline{w} = \lim_{n\to\infty} \tfrac{1}{n} \sum_{k\le n} w_k,$$

whenever the limit exists (the same bound holds for the limit superior otherwise). Since $w_k \le W$, $\overline{w} \le W$, with equality if and only if $w_k \to W$ in density, that is, if and only if the policy switches at the guard thresholds on all but a vanishing fraction of cycles, which is the rule-induced hysteresis on the states it visits. Every other policy therefore has a switching frequency, and so a cost, no smaller. □

Under constant drift the maximal-width cycle attains only the midpoint of $\mathcal{F}$ as its mean; at any other prescribed mean narrower cycles must be used and the argument no longer selects the hysteresis. The effect of exponential drift and of the 15-min step on the sampled cycle mean is measured in Supplementary Section S2; matching in the protocol is done on the simulated trajectory, never on the midpoint formula.

Remark 5 (Room inside the rules) Neither hypothesis of Corollary 1 holds in the protocol. The disturbance is diurnal, and the matched mean lies below the midpoint of the training rule set's free interval $\mathcal{F}$ (23.10 °C for R1/R2 at 23.80/22.40 °C). The guard hysteresis reaches the target only by lowering its own ceiling (fitted R1 23.60–23.66 °C), which narrows $\mathcal{F}$ and leaves the corollary silent about policies in the original $\mathcal{F}$. What room it leaves is measured with causal controllers (Section 8.3). Inside the wider training rule set a 24 h receding-horizon controller saves more than the fitted guard in Riyadh and Jeddah, at a cooler mean; inside the guard's own fitted rules it does not beat the guard by a resolved margin. The D1 shield does not use those rules: it replaces their fixed margins by the plant's exact one-step prediction (Section 5.5), and reaches the tuned thermostat's level on a latch in every city and on the trained policies in Riyadh and Jeddah. Part of the room the rules leave is therefore their fixed margins, which a one-step model removes, and part is reachable by lookahead inside the wider rule set; the comparisons carry the comfort-offset caveat of Section 10.4.

Within the constant-rate abstraction the duty ratio, and so the energy, is the same for every policy (Step 2). Under the exact, temperature-dependent dynamics it is not: away from matched comfort, switching off early raises the mean temperature, which *reduces* energy by (9) while adding a switch, so the trade can go either way. Matching the mean removes the energy term and leaves the switching term, which is why every arm in Section 8 is comfort-matched before its cost is quoted. Under the diurnal drift of the evaluated problem the corollary does not hold, so PPO's coincidence with the guard hysteresis (Section 9.2) does not confirm it.

The corollary is narrow, and its hypotheses locate where a better policy can act: a diurnal disturbance or a time-varying price makes the cost-minimising switch points time-dependent, an occupancy schedule changes the band, and a shared capacity-limited plant couples the zones' action sets.

Without setback, any saving is therefore realised by avoiding on-commands, not by precooling (Remark 2; Section 8.2).

5 Controllers

Section 4.4 establishes that O1 must hold at every interval while O2 and O3 are economic. The rule-based reinforcement learning controller (RBRL) splits the controller along exactly that line: a rule layer that enforces O1 by construction, and a learned layer that works on O2 and O3 inside the region the rules leave free. Its comparators, rule-based, model-based and clairvoyant, are defined in Section 5.5.

5.1 Architecture

The two layers are shown in Figure 2. The rule layer overrides the agent's sampled action whenever the zone temperature reaches a guard threshold. It is applied in the environment after sampling, a post-posed shield (Alshiekh et al. 2018), so the learner stores its own sampled action and treats the override as part of the dynamics; invalid-action masking inside the policy (Huang and Ontañón 2022) would remove the forced steps from the gradient but, with two actions, leaves the optimum unchanged. The PPO layer chooses the action where the rules leave it free; Section 9.2 tests what it achieves there.

5.2 State Space

The numerical study controls a single zone, and at interval t the agent observes eight features:

$$s_t = \Big(T^{(t)},\ u^{(t-1)},\ T^{(t)}_{\text{out}},\ T^{(t+4)}_{\text{out}},\ I^{(t)},\ \sin\tfrac{2\pi h_t}{24},\ \cos\tfrac{2\pi h_t}{24},\ t/N\Big) \in \mathbb{R}^8, \tag{16}$$

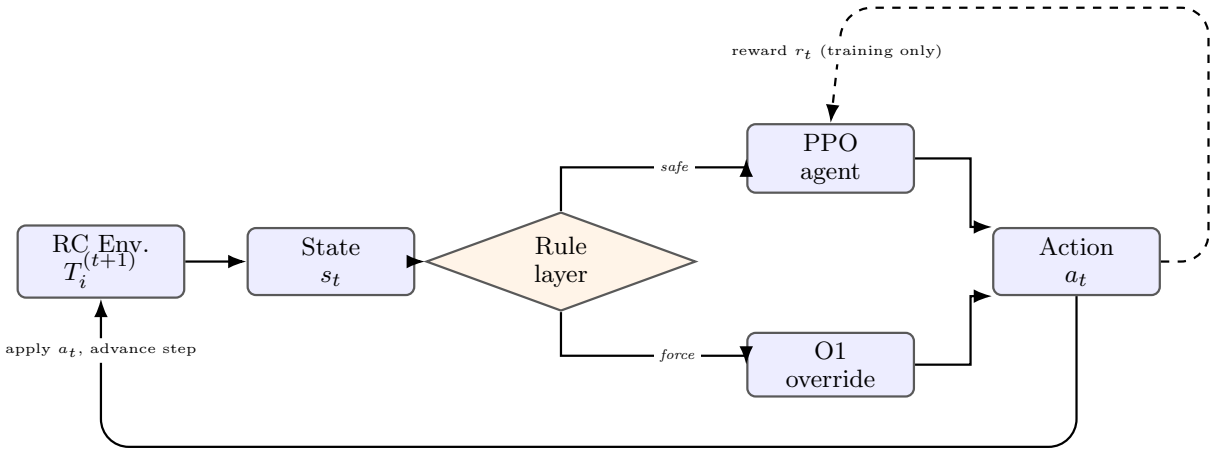


Fig. 2: RBRL architecture. The rule layer routes to the PPO agent when no comfort boundary is threatened (*safe*) and overrides it with the O1-enforcing action when one is imminent (*force*); the dashed reward loop is used in training only.

namely the zone temperature, the previous action, the outdoor temperature now and one hour ahead, the south-vertical irradiance, the hour of day as a sine–cosine pair, and the fraction t/N of the episode elapsed, each scaled to order one (scalings in the repository). The one-hour outdoor value is a perfect forecast. The effective inter-zone temperature $T_{\mathrm{iz},i}$ (equation (6)) is not observed, because there is no neighbouring zone, and neither is E_{cum}: under the in-force two-block tariff this dwelling never leaves the first block (Remark 3), so the cost rate does not depend on it. A learned multi-zone or stepped-tariff policy would need both; none is trained here.

5.3 Rules and Action Space

Definition 7 (Rule-based overrides)

$$\textbf{R1 (hard, force on:)}\ T_i^{(t)} \geq T_{\max} - \epsilon_{\mathrm{hot}} \ \Rightarrow\ u_i^{(t)} \leftarrow 1. \tag{17}$$

$$\textbf{R2 (hard, force off:)}\ T_i^{(t)} \leq T_{\min} + \epsilon_{\mathrm{cold}} \ \Rightarrow\ u_i^{(t)} \leftarrow 0. \tag{18}$$

The guard bands are *asymmetric* and are derived, not assumed. Each must absorb the temperature excursion of one actuation interval in its own direction:

$$\begin{aligned} \epsilon_{\mathrm{cold}} &\geq \frac{Q_{\mathrm{hvac}}\Delta t}{C_{\mathrm{th}}}, \\ \epsilon_{\mathrm{hot}} &\geq \frac{\Delta t}{C_{\mathrm{th}}} \max_t \left[\Lambda_i (T_{\mathrm{out}}^{(t)} - T_{\max}) + Q_{\mathrm{sol}}^{(i,t)}\right]. \end{aligned} \tag{19}$$

The two differ because cooling is driven by the plant and warming by the load. At the parameters of Table 2 the cooling excursion is 0.375 °C; the peak seasonal load of 690, 664 and 513 W across the three cities gives warming excursions of 0.129, 0.125 and 0.096 °C. Cooling is about 2.9 times the fastest warming, so $\epsilon_{\mathrm{cold}} = 0.40$ and $\epsilon_{\mathrm{hot}} = 0.20$ °C are adopted, so the agent acts on $(22.40, 23.80)$ rather than the $(22.5, 23.5)$ a symmetric 0.5 °C guard would leave. Even so, the rule layer can exclude the tuned thermostat's own switch points (23.71/22.36, 23.70/22.34 and 23.64/22.32 °C in Riyadh, Jeddah and Taif, all with an off point below R2): that exclusion is a property of the method, since any fixed guard band forbids some members of the family it guards.

The rule layer contains no pre-cooling rule, because Proposition 1 and Remark 2 rule it out under every cost model used here. Pre-cooling would pay only under a physical condition, not a tariff one: shifting cooling from t_2 back to t_1 pays only if the price of *delivered* cooling falls by more than the store leaks over the interval, which at $\tau = C_{\mathrm{th}}/\Lambda_i = 56.1$ h is 3.5% over a four-hour shift and 6.8% over eight, averaged over the shift;

the end-point figures $1 - e^{-\Delta t/\tau}$ are 6.9% and 13.3%. A cumulative-block tariff supplies none of that. A temperature-dependent coefficient of performance would supply 15–20% across a typical diurnal swing and is the route by which a pre-cooling rule could be reinstated (Section 10.4). R1 and R2 enforce O1 by construction, not by training. The action is $a_t = (u_1^{(t)}, \ldots, u_{N_z}^{(t)}) \in \{0,1\}^{N_z}$. A single PPO head is used.

5.4 Reward Function and PPO Architecture

The per-step reward is the negative interval cost minus a linear degree-hour comfort penalty:

$$r_t = -c^{(t)} - \alpha \left[\max\!\left(0, T^{(t+1)} - T_{\max}\right) + \max\!\left(0, T_{\min} - T^{(t+1)}\right) \right] \Delta t, \tag{20}$$

with $\alpha = 50\,\mathrm{SAR\ per\ ^\circ C\,h}$, and at the last step of an episode the terminal thermal state is priced exactly as the dynamic programme and the simulator price it. Under the rule layer the penalty is zero except on plant-limited cold drift. The policy and value networks are separate, each with two fully connected layers (256–128 units, tanh); the policy outputs the logits of a categorical distribution over $u \in \{0,1\}$. Training uses Stable-Baselines3 v2.9.2 (Raffin et al. 2021) with Adam at a constant learning rate of 3×10^{-4}, $\gamma = 0.999$, clip coefficient 0.2, entropy coefficient 0, rollouts of 1,344 steps (two episodes) with mini-batches of 336, and otherwise default settings; no hyperparameter search was made. Training runs for 2×10^5 environment steps (about 298 episodes of 672 steps, one 168 h window at $\Delta t = 15\,\mathrm{min}$) with five seeds (42–46), none discarded, a fixed budget and no checkpoint selection; each episode draws its start date from the six tuning windows of Section 7. During training the ceiling is the band edge, $\overline{T} = T_{\max}$, so R1 fires at 23.80 °C and R2 at 22.40 °C.

No convergence claim is made. With the rules on and a flat price, the training reward, which has no term in the mean temperature, is nearly minimised by holding the previous action until a rule fires (within 0.4, 1.4 and 2.3% of the exact $\lambda = 0$ dynamic-programming optimum inside the rules, called E1; Section 9.2). In Riyadh and Jeddah the seeds learn that policy: eight of ten runs reproduce the guard hysteresis's pooled saving to 10^{-5}, and the other two differ by less than 0.004 points; in Taif four of five seeds fall below it. That shortfall is one of generalisation, not of training length: scored with the training cost at the training rules, three of those four seeds are within 0.05% of the hold policy on the six tuning weeks (the fourth 1.4%) and 0.5–2.5% above it only on the held-out weeks, so a longer run on the same weeks has almost nothing left to gain. A second variant, E2, tests whether this is a property of the reward: it adds the stage term $(\lambda / n_{\mathrm{occ}}) \max(T^{(t+1)}, T_{\min})$ of the receding-horizon controller, with λ taken from the in-rules 6 h controller, and a perfect outdoor forecast at +1, +2, +3 and +6 h and irradiance at +1 and +3 h to the state; it is trained in Riyadh with three seeds, an entropy coefficient of 0.01 and reward normalisation, under the same rules in training and evaluation, with no Rule A selection. It stays close to the rules' hold policy (Section 9.2). Algorithm 1 sets out training and schedule extraction together; the policy π_θ it returns is the one evaluated in every RBRL row of Section 8.

Algorithm 1 RBRL: training, ceiling selection and evaluation

Require: simulator (4); tuning windows $\mathcal{W}_{\text{tune}}$ and held-out windows $\mathcal{W}_{\text{eval}}$ (start dates d_0, $N = 672$ steps each); price p; c_{sw}; budget $B = 2 \times 10^5$ steps; guard margins $\epsilon_{\text{hot}}, \epsilon_{\text{cold}}$; ceiling grid $\mathcal{C} = \{23.50, 23.51, \ldots, 24.00\}\,^\circ\text{C}$

Ensure: policy π_θ, ceiling $\overline{T}^*$, held-out costs $\{J_w\}_{w \in \mathcal{W}_{\text{eval}}}$

Phase 1: training ($\overline{T} = T_{\max}$)

Initialise π_θ; $n \leftarrow 0$
while $n < B$ **do** ▷ fixed budget; no early stopping
  Draw $d_0 \in \mathcal{W}_{\text{tune}}$; $T^{(0)} \leftarrow 23\,^\circ\text{C}$, $u^{(-1)} \leftarrow 0$
  for $t = 0$ **to** $N - 1$ **do**
    observe s_t (16); $a_t \sim \pi_\theta(\cdot \mid s_t)$
    $u^{(t)} \leftarrow \mathcal{R}(a_t, T^{(t)}; \overline{T})$ ▷ R1–R2, eq. (19)
    step (4); r_t by (20); $n \leftarrow n + 1$
  end for
  add the terminal price of $T^{(N)}$ to r_{N-1}
  every 1,344 steps: PPO update (clipped surrogate, GAE)
end while

Phase 2: R1-ceiling selection on the tuning windows (Rule A)

for $\overline{T} \in \mathcal{C}$ **do**
  run $\arg\max_a \pi_\theta$ with $\mathcal{R}(\cdot; \overline{T})$ on every $w \in \mathcal{W}_{\text{tune}}$
end for
$\overline{T}^* \leftarrow$ admissible $\overline{T}$ (Rule A) of least equalised tuning cost ratio

Phase 3: held-out evaluation (frozen π_θ, $\overline{T}^*$)

for $w \in \mathcal{W}_{\text{eval}}$ **do**
  run $\arg\max_a \pi_\theta$ with $\mathcal{R}(\cdot; \overline{T}^*)$; record J_w and $\tilde{\bar{T}}_w$
end for
return π_θ, $\overline{T}^*$, $\{J_w\}$ ▷ compared by Rule C, Section 7

5.5 Competing Methods

THERM (unoptimised baseline) models the automatic remote-controller thermostat common in Saudi residential buildings (Alotaibi et al. 2024): on when $T_i \geq T_{\max} - 0.5$, off when $T_i \leq T_{\min} + 0.5$, with no tariff awareness. On the occupied band adopted here this is a 1 °C deadband centred on 23 °C, assumed, not measured, to represent the split units in this stock (Section 9.2). Every optimising controller gains over THERM by using more of the 2 °C band under the same comfort obligation.

TUNED (rule-based comparator) is the same hysteresis family with its two switch points free. They are searched on a 0.05 °C grid and then refined iteratively on a 0.01 °C grid around the best admissible pair; a row at the 0.1 °C resolution of typical hardware is in the supplement. Without TUNED, any gap attributed to learning or planning may in fact be obtainable by commissioning the thermostat properly.

Guard hysteresis is the hysteresis at the rule layer's own thresholds, R1 and R2, with no learned policy: on at $\overline{T} - \epsilon_{\text{hot}}$, off at $T_{\min} + \epsilon_{\text{cold}}$, with the ceiling $\overline{T}$ fitted like every other parameter. **D1** is a one-step model shield on a latch that holds the previous action and has no R1/R2 guard bands: it forces the unit on when the exact one-step prediction with the unit off would end above $\overline{T} - \sigma$, and off when the prediction with it on would end below $T_{\min} + \sigma$, with its own ceiling $\overline{T} \leq T_{\max}$ fitted for each margin $\sigma \in \{0, 0.05, 0.10\}$ K; on the trained PPO policies it replaces the rule layer. It needs the plant model and the current outdoor temperature and irradiance, where TUNED needs only the zone temperature. D1 separates what an exact one-step model buys from what a longer horizon buys.

MPC (causal receding-horizon comparator, also "the planner"). At every step the controller solves, by dynamic programming over a temperature grid with the first step taken from the exact state, the horizon problem with stage cost $c^{(t)} + (\lambda / n_{\text{occ}}) \max(T^{(t+1)}, T_{\min})$ over occupied steps, applies the first action and re-solves. Cooling steps that would end below $T_{\min}$ are infeasible, as for every optimiser here. The multiplier λ is the comfort device: one value per city, horizon and forecast mode, fitted on the tuning windows by Rule A of Section 7 and then frozen. If no multiplier is admissible even at the 0.10 K fallback, the planner's occupied ceiling is lowered by a back-off b fitted by the same rule (cells marked †); the order is λ at 0.03 K, then λ at 0.10 K, then b. For the 6 and 24 h perfect-forecast controllers the fitted values lie within 2.5% of $\lambda^* = \Lambda_i\,(168\,\text{h})\,p/\text{COP}$, the marginal envelope cost of 1 K over a window,

at which the Lagrangian is indifferent to the temperature level; the controller therefore minimises on-commands and its comfort level is set by the Rule A selection of λ. Close to λ^* the result is sensitive to λ: in Riyadh, neighbouring admissible multipliers give 24 h savings that differ by 1.4–1.8 points at both grid resolutions, so the fitted λ/λ^* is reported in Section 8.3. The horizon problem is closed by a rollout terminal cost: the cost, including the on-command charge and the λ term, of running TUNED for 24 h from the terminal state on a persistence forecast, plus the price of the stored heat. This is a truncated rollout of a base policy (Bertsekas 2020); the classical guarantee of closed-loop cost no greater than the base policy's does not carry over, because the tail is truncated, uses persistence weather and is evaluated on a grid, so any improvement over TUNED is an empirical finding. Horizons of 0.25, 1, 2, 6, 24 and 48 h are run with two forecast modes: *perfect*, the true in-horizon weather, and *persistence*, in which the forecast for step j is the observation one day earlier, with the current outdoor temperature measured. In both modes the controller plans on the plant's own model with the exact state. The *in-rules* variant restricts every action to the training rule set, R1/R2 at 23.80/22.40 °C; its terminal rollout still uses TUNED, whose off point lies below R2. The value grid has 1,601 nodes; at the fitted λ, refining it to 3,201 nodes changes the 24 h saving by 0.10, 0.45 and 0.29 points in Riyadh, Jeddah and Taif; the grid uncertainty is therefore taken as at most 0.5 points, below the 3-point resolution rule.

CLV (clairvoyant reference) is a dynamic programme over the exact dynamics (3) that sees the whole window's weather. Its multiplier is set per held-out window, by bisection (with an occupied-ceiling cap as a second knob where needed), aiming at a clipped mean within 0.002 K of THERM's. At the critical multiplier the optimal schedule's mean jumps by 0.1–0.3 K, so the target is met in only 2 of 7 windows in Riyadh, 1 of 7 in Jeddah and none in Taif; elsewhere the closest schedule no warmer than THERM is kept, and its residual is recorded (per window in the supplement).

End-of-window convention. The simulator, and so every causal arm, prices the terminal thermal state at the electricity needed to return it to the initial temperature T_0 at rated COP. A programme that knows where the window ends could exploit that price, because it omits the on-command the return would need. The oracles (CLV, the bound and B_a) therefore use $V_N = \left[p\,C_{\mathrm{th}}/(3.6\times10^6\ \mathrm{COP})+c_{\mathrm{sw}}/(1\ \mathrm{K})\right](T_N-T_0)$: heat left above T_0 is charged its re-cooling energy plus the fraction $(T_N - T_0)/1\,\mathrm{K}$ of one on-command, that is, one on-command per kelvin, which is a convention. The savings of the causal arms in Table 3 use the simulator terminal, the bracket and B_a use V_N, and B^* is also reported under the simulator terminal for Riyadh and Jeddah. E1 keeps the simulator terminal, which is PPO's own objective. CLV and the bound of Section 8.1 bracket what a schedule no warmer than the baseline could save; the gap between TUNED and that bracket is the headroom any controller could in principle occupy.

Genetic-algorithm and simulated-annealing schedulers (El Makroum et al. 2023; Bandarra Filho et al. 2025) are not run. Both are day-ahead offline solvers, and the bound of (23) applies to them as to every other schedule at the baseline's mean, so it answers the question they would be introduced to answer. No third-party learned comparator (SAC, DQN) is run inside the rule layer either. A relaxed bound restricted to the training rule set (14.9, 15.6 and 16.2% in Riyadh, Jeddah and Taif) caps what any such learner could attain, but whether another learner would come closer to it than PPO is untested, as is any learner acting outside R1–R2.

The framework is applied to the three climates characterised next.

6 Case Studies: Riyadh, Jeddah and Taif

Three sites are used, chosen to span the conditions under which a Saudi dwelling is cooled, not to represent the Kingdom: hot-arid inland (Riyadh), hot-humid coastal (Jeddah) and mild highland (Taif).

6.1 Climate Profiles

Riyadh (BWh, 24.7°N, 635 m) is an inland desert city with summer highs reaching 46.6 °C and a cooling-season diurnal swing of 12.7 K, the largest of the three sites (CDD_{18} = 3,693; cf. (Haddad et al. 2024)). Jeddah (BSh, 21.7°N, 15 m)

is a Red Sea coastal city, classified hot-humid (Rodrigues et al. 2025); in the record used here its annual mean relative humidity is 58.5% with 17.5% of hours above 70%, and it has a compressed diurnal swing of 8.4 K and little overnight relief (CDD_{18} = 4,260; the ordering of the two totals depends on the typical-year record, and what distinguishes the sites here is the diurnal swing). Taif (BSh, 21.5°N, 1,478 m) is a highland site whose altitude lowers both the peak and the mean: over the cooling season its mean outdoor temperature is 28.6 °C against Riyadh's 35.0 and Jeddah's 32.7; its peak 39.4 °C against 46.6 and 44.0; and its CDD_{18} = 2,277 against 3,693 and 4,260. Its diurnal swing, 11.7 K, is close to Riyadh's. The load is therefore the smallest of the three while the headroom is not, which is why Taif leaves the least to schedule in absolute terms; the lower air density at that altitude is also why the infiltration conductance is evaluated per city (Section 6.2); monthly and diurnal profiles are in the repository. All three simulations use TMYx typical-year records from `climate.onebuilding.org`, WMO 404380 (Riyadh, period of record 2009–2022, 12 years of data), 410240 (Jeddah, 2009–2023, 15 years) and 410360 (Taif, 2009–2023, 15 years); hourly dry-bulb temperature and the direct-normal and diffuse-horizontal irradiances are read from fields 7, 15 and 16, and the glazing gain is the plane-of-array irradiance on a south-facing vertical surface under an isotropic sky with ground albedo 0.30.

6.2 Building Code, Tariff, Humidity, and Comfort Standard

SBC 2018 mandates opaque-wall U-value $\leq 0.45\,\mathrm{W\,m^{-2}K^{-1}}$ and window SHGC ≤ 0.25 in both Climate Zone A (Jeddah) and Zone B (Riyadh) (Alayed et al. 2021). Table 2 meets both at assembly level: $U_{\mathrm{wall}} = 0.45$ and SHGC $= 0.25$ exactly, with $U_{\mathrm{win}} = 2.8\,\mathrm{W\,m^{-2}K^{-1}}$ consistent with double glazing. The multi-tier stepwise tariff follows equation (11); for Models L and E, $c_r^{\mathrm{L}} = 0.12$ SAR/kWh is adopted as an illustrative flat rate, below the post-2018 first-block price of 0.18 SAR/kWh. It is *not* the consumption-weighted average of (11) over this account, which at ≈884 kWh is the first-block rate of 0.05 SAR/kWh exactly. The occupied comfort band $[22, 24]$ °C is consistent with the *thermal* criterion of ASHRAE 55-2023 (ASHRAE 2023) for mechanically conditioned residential spaces (the standard's humidity criterion is not evaluated here; Section 10.4). Bienvenido-Huertas et al. (2024) validate this use, and its upper limit reflects Saudi practice: Krarti (2020) recommends 23–25 °C for Kingdom dwellings and observes setpoints as low as 18 °C in the stock. Lamsal et al. (2023) review adaptive thermal comfort models and find wider deadbands appropriate for energy-efficient design in hot climates; the narrower band adopted here raises the reported saving through the baseline's cycling (Section 10.3) and was not swept. A fixed band is adopted because Saudi residential buildings are predominantly mechanically cooled; ASHRAE 55 Section 5.3 restricts adaptive comfort to naturally ventilated spaces. The saving does not transfer across band widths, because the actuation condition (8) and the baseline's cycling rate both depend on the width.

7 Simulation Setup

The controllers of Section 5 are evaluated on the sites of Section 6 under one protocol, set out here. All simulations run in Python 3.13.18 with NumPy 2.5.3, integrating equation (3) exactly at $\Delta t = 15$ min, the interval that satisfies (8) with a factor of 2.7 (Remark 1). Solar gain is $Q_{\mathrm{sol}}^{(i,t)} = \mathrm{SHGC} \times A_{gl} \times I^{(t)}$, with $I^{(t)}$ the plane-of-array irradiance of Section 6.1; hourly weather values are held constant over the intervals of each hour, without interpolation. The RC model has no occupancy or internal gains.

Background load and tariff.

The tariff is levied on the whole account, so the cumulative energy E_{cum} starts from a non-HVAC load of $E_{\mathrm{base}} \approx 800$ kWh/month (Al-Homoud and Krarti 2021). The single zone adds about 84 kWh a month, so the account reaches ≈884 kWh and crosses no boundary of any schedule; sixteen such zones would approach the first boundary of the historical four-tier schedule, which is why Model S is kept as a stress test (Remark 3).

Evaluation protocol.

Evaluation is confined to the cooling season, 1 May to 31 October, because the plant is cooling-only. Thirteen windows of 168 h start every 14 days, at

zero-based day indices 121 to 289 (1 January is day 0), none overlapping, and are assigned alternately within one TMYx year. The seven *held-out* windows start on 2 May, 30 May, 27 June, 25 July, 22 August, 19 September and 17 October; the six *tuning* windows start on 16 May, 13 June, 11 July, 8 August, 5 September and 3 October. There is no separate validation set: every parameter, including PPO's ceiling, is selected on the tuning windows alone, PPO trains for a fixed budget with no early stopping or checkpoint selection, and the held-out windows are touched only once, for evaluation. The seven held-out windows span the cooling season, so the seasonal variation is carried by the per-window tests rather than by a separate seasonal evaluation; the inter-annual spread is measured on actual years (Section 10.2). One protocol, implemented in one code module, applies to every arm and oracle.

Comfort. On every occupied step the post-step temperature must not exceed $T_{\max}$; the occupied over-temperature in degree-hours is recorded for every arm. On the cold side, *cooling may never end a step below $T_{\min}$; free-float is unrestricted*: $u^{(t)} = 1$ implies $T^{(t+1)} \geq T_{\min}$ at every step, and $u^{(t)} = 0$ has no lower limit. Sub-band states can therefore arise only by free drift when the off-state equilibrium lies below the band. In the optimisers this is a hard constraint on cooling transitions. The rule-layer arms and D1 are designed to meet it from the nominal plant parameters, and TUNED meets it on the tuning windows by Rule A and threshold placement, not by guarantee; for all of them it is measured as commanded under-temperature, the degree-hours below $T_{\min}$ at the end of cooling steps, reported beside the free-float under-temperature of idle occupied steps.

Rule A (selection on the tuning windows). Each causal arm has one parameter set per city: TUNED's two switch points, the guard's and RBRL's ceiling, D1's margin σ, and the receding-horizon controller's λ for each horizon and forecast mode. A candidate is admissible if, on every tuning window, it has zero over-temperature and zero commanded under-temperature, and if its clipped mean occupied temperature pooled over the tuning windows lies within 0.03 K of THERM's. Among admissible candidates the one with the least equalised tuning cost ratio $\sum_w J_w / \sum_w J_{\mathrm{THERM}(\delta_w),w}$ (Rule C applied to the tuning windows) is selected and frozen. If no candidate meets the 0.03 K condition, it is relaxed to 0.10 K, for every arm alike; each table states which applies (`fit_status`), and an arm with no admissible parameter is reported as such and not tabulated. A sensitivity refit under a per-window 0.03 K criterion is in the supplement.

Rule B (clipped mean). Occupied means are computed on $\tilde{T} = \max(T, T_{\min})$ for every arm and for THERM, so a free-float cold excursion buys no credit against warmth elsewhere.

Rule C (equalisation per held-out window). On each held-out window w the frozen arm is compared with THERM shifted by δ_w (on at $23.5+\delta_w$, off at $22.5+\delta_w$), with δ_w scanned over $[-0.4, 0.4]$ K in 0.002 K steps. THERM's clipped mean is not monotone in δ, so the arm's mean can be reached at several shifts; the comparator cost is the *minimum* over all crossings of the piecewise-linear scan, interpolated linearly within each bracketing segment, which credits the baseline with its cheapest setting of equal comfort and can only understate a saving. Scored this way, the unshifted thermostat itself saves -0.8, -0.1 and 0.0% in Riyadh, Jeddah and Taif, which is the size of the rule's conservatism; the range between the cheapest and the dearest crossing is in the supplement. The only saving printed is the pooled

$$s = 1 - \frac{\sum_w J_{\mathrm{arm},w}}{\sum_w J_{\mathrm{THERM}(\delta_w),w}}, \tag{21}$$

over the seven held-out windows; the raw saving against unshifted THERM, the pooled and largest per-window mean differences, the comfort metrics and the pooled end temperature are printed beside it. Per-window savings enter only the tests. No held-out window contributes to any arm's parameters. CLV is matched per held-out window (Section 5.5), and the bound needs no selection (Section 8.1).

Statistics. The unit is the per-window difference in equalised saving between two arms on the same seven windows. Each contrast is tested with the exact two-sided Wilcoxon signed-rank and sign tests and summarised by the Hodges–Lehmann estimate with an exact 95.3% interval (the third smallest and largest of the 28 Walsh averages); the estimate and interval lead, the p-value is secondary. The primary families, each one contrast in three cities with Holm's correction within

the family, are TUNED against THERM, MPC (24 h, perfect) against TUNED, MPC (24 h, persistence) against TUNED, and the guard against TUNED; E2 is not tested against the guard, because its comfort is not matched to the other arms. With $n = 7$ and three cities, 0.047 is the smallest Holm-adjusted p attainable, reached by a city only when all seven of its windows agree in sign; one discordant window removes significance, and such results are reported as counts. The horizon, capacity, switching-charge, deadband and tariff sweeps are descriptive. Pooled differences of less than 3 points between arms are treated as not resolved, whatever a test says, because the occupied-mean staircase and the tuning noise are of that order; the threshold does not absorb comfort offsets between arms (Section 10.4). The held-out weeks come from one TMYx year, so the tests describe those weeks. Estimates, intervals, p-values and comfort offsets for every family are in Supplementary Table S8.

Switching charge. For a split unit of ∼2,500 SAR installed cost, a start life of order 10^5 gives 0.025 SAR per start, and a 15-year life at about 1,460 on-commands a year (four a day all year) gives about 0.114 SAR; counting only the 184-day cooling season at THERM's 4.10 a day would give about 0.22 SAR. The range 0.025–0.114 SAR per on-command is therefore a convention, not a measured wear cost; none was found for Saudi split units (Pang and Liu 2011). The tables use $c_{\mathrm{sw}} = 0.15$ SAR as the computational reference, because the PPO policies were trained at it; it lies above the derived range and overstates the saving against it, since the saving rises with c_{sw}. Results at 0.025 and 0.114 SAR are reported beside it (Section 9.2). The cost figures are therefore a pricing of the on-command reduction, which is the quantity actually measured.

8 Results

Results are reported on the seven held-out windows at the tariff in force, $c_{\mathrm{sw}} = 0.15$ SAR, $\Delta t = 15$ min and $C_{\mathrm{th}} = 4{,}800$ kJ/K, with the receding-horizon controller on a 1,601-node grid. Every fitted single-zone arm records zero occupied over-temperature on the held-out windows; the two fixed setback arms do not (Section 8.2). Commanded under-temperature (Supplementary Table S4) is zero for every optimiser, rule-layer and D1 arm except the Riyadh 6 h persistence controller, whose plan meets the cold rule on the forecast but not on the realised weather (0.0006 °C·h over the seven windows). The multi-zone controller, which plans each room against its neighbours' plans of the previous step, records 0.0001 °C·h of occupied over-temperature in the 1×4 row at $\lambda_{ij} = 35$ W/K and at most 0.00003 °C·h of commanded under-temperature in the 2×2 grid. The threshold arms meet it only on the tuning windows: on the held-out windows TUNED records 0.0007 and 0.0066 °C·h in Jeddah and Taif, and TUNED-SB 0.0102 °C·h in Riyadh, the largest in the headline configuration (0.024 °C·h at most in any sensitivity run). The largest free-float value, 0.96 °C·h, belongs to the Taif 24 h persistence controller, and in Taif TUNED and the guard also free-float below the band (0.29 and 0.34 °C·h). Section 8.1 first fixes what clairvoyance is worth; Sections 8.2–8.4 then read the controllers against it.

8.1 What Clairvoyance Is Worth: a Bound and an Attained Schedule

The bound keeps comfort pointwise and applies the comfort match only in the mean, so it covers every schedule no warmer than the baseline:

$$\begin{aligned} \min_{u} J(u) \quad \text{s.t.} \quad & T^{(t+1)} \le T^{\mathrm{hi}}_t \ \ \forall t, \\ & u_t = 1 \ \Rightarrow \ T^{(t+1)} \ge T_{\min} \ \ \forall t, \\ & \tfrac{1}{|\mathcal{O}|} \textstyle\sum_{t \in \mathcal{O}} \tilde{T}^{(t+1)} \ \le \ T^{\star}, \qquad \tilde{T} = \max(T, T_{\min}), \end{aligned} \tag{22}$$

with $\mathcal{O}$ the occupied steps, $T^{\mathrm{hi}}_t = T_{\max}$ while occupied and the empty-house limit otherwise, and $T^{\star}$ the baseline's clipped mean. A cooling step may never end below the band; a free-floating step may, because a cooling-only plant cannot prevent it, and the clipped mean gives such an excursion no credit. The mean constraint is an inequality because only warmth can buy a saving, and because with binary actions on a 15 min grid the attainable means form a finite set. Relaxing it alone gives $L(u, \lambda) = J(u) + \lambda\big(|\mathcal{O}|^{-1} \sum_{t \in \mathcal{O}} \tilde{T}^{(t+1)} - T^{\star}\big)$, the same dynamic programme with one additive stage cost, and for

any $\lambda \geq 0$ and any u feasible for (22)

$$J(u) \;\geq\; L(u,\lambda) \;\geq\; \min_{u'} L(u',\lambda) \;=\; L^{\star}(\lambda). \quad (23)$$

The attained mean is a step function of λ, so the dual gap need not close and no exactly mean-matched schedule need exist in the Lagrangian family.

A provable bound on a grid. $L^{\star}(\lambda)$ is computed by dynamic programming on a uniform temperature grid of spacing h, and an ordinary discretised programme can return a minimum that is too *expensive*, the unsafe direction for a bound. The reported bound therefore uses an optimistic-neighbour relaxation. For each node y, previous action and action u, the exact successor $f_u(y)$ is computed; energy and on-command costs are charged exactly; the mean term is charged at $\max(f_u(y) - h, T_{\min})$; feasibility is tested against limits widened by h; the continuation is the smaller of the values at the two nodes bracketing $f_u(y)$; the terminal value is evaluated at $y - h$; and the initial value is the smaller of the values at the two nodes bracketing $T^{(0)}$. The grid covers the band plus 0.1 K and extends below it to the window's lowest free-float equilibrium, so no trajectory admitted by (22) leaves it.

Proof. Take any feasible u with true trajectory x_t. Let y_0 be the node bracketing x_0 nearest to it, and y_{t+1} the node bracketing $f_{u_t}(y_t)$ nearest to x_{t+1}. The plant map $f(x) = T_\infty + (x - T_\infty)a$, $a = e^{-\Delta t \Lambda / C_{\text{th}}} \in (0,1)$, gives $|f(y_t) - x_{t+1}| = a|y_t - x_t|$, so $|y_t - x_t| \leq h$ implies $|y_{t+1} - x_{t+1}| \leq h$, and by induction this holds for every t. Along this path each relaxed stage cost is at most the true one: the mean term because $f(y_t) - h \leq f(x_t)$ and $\lambda \geq 0$; the terminal value because it increases with T and $y_N - h \leq x_N$. Every true-feasible step is relaxed-feasible, because $|f(y_t) - f(x_t)| \leq h$ (and, for the rule-constrained variant, whose masks test the current state, because $|y_t - x_t| \leq h$). Because the relaxed continuation takes the smaller value of the bracketing pair, the relaxed value at y_0 is no larger than the relaxed cost along this path. Hence $L_h^{\star}(\lambda) \leq J(u) + \lambda\big(\overline{\tilde{T}}(u) - T^{\star}\big) \leq J(u)$ for every $\lambda \geq 0$ and every u feasible for (22), pooled over the held-out windows. $L_h^{\star}$ is concave in λ, so λ is maximised on an eighteen-point grid with ternary-search refinement; an under-maximised λ weakens the bound without invalidating it. □

The relaxation has a price. The bracketing minimum may round the successor down by up to h at every step, and the roundings accumulate over the plant's memory, $1/(1-a) \approx 225$ steps, so the slack is of order $h/(1-a)$ in temperature and shrinks roughly in proportion to h: under V_N the relaxed saving bound is 21.3, 20.6, 20.2 and 20.0% in Riyadh at 1,601, 3,201, 6,401 and 12,801 nodes (21.2, 20.5, 20.1 and 19.8% in Jeddah), against 19.5% and 19.1–19.4% from the ordinary programme, which is not a bound; at 12,801 nodes the slack is 0.47 and 0.42 points. In Taif the relaxed bound is 23.8, 22.2, 21.4 and 20.9% against 19.9–20.2%, a slack of 0.75 points at 12,801 nodes. Every grid gives a valid bound, so the smallest is reported, rounded up to 0.1. The bound is a minimum over every action sequence in the window under the same cost, so it holds for any controller whatever information it uses; the receding-horizon controller's forecast beyond the window end cannot take it above the bound, because post-window weather enters the cost only through the terminal price of the stored heat.

What is reported. The headline bound is $B^* = 1 - L^{\star}(T^{\star}) / \sum_w J_{\text{THERM}(0),w}$ at THERM's pooled mean. Each arm a is warmer or cooler than THERM by its own $\Delta\tilde{T}$, so its bound is evaluated at its own mean, $B_a = 1 - L^{\star}(\overline{\tilde{T}}_a) / \sum_w J_{\text{THERM}(\delta_w^a),w}$, and $s_a \leq B_a$ then holds exactly once each arm is rescored under V_N; B_a for the principal arms is in Supplementary Table S5 (all 89 rows in the results files). The lower end of the bracket is the *best attained schedule no warmer than THERM*, scored as $1 - \sum_w J_{\text{CLV},w} / \sum_w J_{\text{THERM}(0),w}$, with the same denominator and the same mean constraint as B^*; the upper end is the provable bound under the model; the gap between them is the sum of the duality gap, the relaxation slack and the replay's own sub-optimality. CLV's equalised saving under Rule C is printed only beside B_{CLV}, the bound at CLV's own mean, never against B^*.

Here the bracket is [18.0, 20.0]% in Riyadh, [18.4, 19.9]% in Jeddah and [16.7, 21.0]% in Taif, 2.0, 1.5 and 4.3 points wide. Taif's relaxation slack (0.75 points) explains little of its width; the rest is duality gap and replay sub-optimality, and CLV meets the comfort target in none of Taif's windows. In an earlier run of the bound under the simulator's energy-only terminal B^*

was 20.0 and 20.1% in Riyadh and Jeddah (not repeated for Taif). CLV's equalised saving is 18.3% against $B_{\mathrm{CLV}} = 19.6\%$ in Riyadh, 18.5% against 19.4% in Jeddah and 16.5% against 20.2% in Taif; end-of-window temperatures are in Supplementary Table S5. Every arm satisfies $s_a \leq B_a$ and every shifted per-window certificate is non-negative under V_N, in all 89 arm and seed rows checked (every causal arm, E2, the D1 shield, the receding-horizon controller in both rule sets, and the setback arms against the setback bound; the rule-layer arms also against the ruled bound). It is a bracket at the instance evaluated (the tariff in force, $c_{\mathrm{sw}} = 0.15\,\mathrm{SAR}$, $\delta_{\mathrm{db}} = 1\,^{\circ}\mathrm{C}$, $\Delta t = 15\,\mathrm{min}$), not at other deadbands or tariffs.

8.2 Cost, Energy and On-Commands

Table 3 gives every arm under the protocol of Section 7, and Table 4 reads the principal arms in three layers. Each saving is the pooled equalised saving against THERM shifted to the arm's own clipped mean, which removes the first-order energy effect of the mean temperature; a residual dependence remains (comfort-offset caveat, Section 10.4). Every parameter was fitted on the tuning windows, so every figure is out-of-sample.

Three layers. Every result is read in three layers (Table 4). The physical headline is the share of on-commands avoided against the comfort-equalised thermostat, which needs no price: 22.8, 24.2 and 21.7% for TUNED, 29.8–33.3% for the 24 h controller with a perfect forecast, and 29.2–42.3% with occupancy setback; Fig. 3 shows where they come from. Beside it, electricity is about unchanged: within ±1.2% for every arm without setback and up to +3.3% for the setback planner, which cools the empty house before arrival. At equal mean temperature every arm uses the same terminal-priced energy (the identity holds to within 4×10^{-5} relative), so the money saved is the charge on avoided on-commands. Priced at the derived 0.025–0.114 SAR per on-command, TUNED is worth 4–19, 5–19 and 2–10 SAR per room per cooling season, and the 24 h controller 6–28, 6–26 and 4–15 SAR; at the 0.15 SAR computational reference, above the derived range, the figures are 25, 25 and 14 and 37, 34 and 20 SAR. The cost difference is therefore a charge on avoided on-commands, assumed to stand for wear, not a lower electricity bill. The pooled cost saving s remains the protocol's ranking metric (Table 3).

The tuned thermostat. Commissioning the two switch points avoids 22.8, 24.2 and 21.7% of on-commands and saves 12.5, 13.6 and 12.4% of cost in Riyadh, Jeddah and Taif, in every held-out window of every city (family F1: Hodges–Lehmann +12.6, +13.7 and +11.8 points, Holm-adjusted $p = 0.047$ in each city), with no model, no forecast and no state beyond the zone temperature. The mechanism is a wider deadband: the fitted switch points (Section 5.3) widen THERM's 1 °C deadband to 1.35, 1.36 and 1.32 °C, close to the widest that the band and the 15-min excursion admit (Remark 1), and the share of on-commands avoided is roughly $1 - \delta_{\mathrm{THERM}}/\delta_{\mathrm{TUNED}}$. The saving is therefore a statement about the comparator's assumed 1 °C deadband as much as about tuning (Section 9.2). At the charges the wear derivation supports, 0.025 and 0.114 SAR, it saves 3.8–11.0%, 4.3–12.0% and 3.9–10.9%. These switch points are fitted on a 0.01 °C grid. A real thermostat sets its thresholds in steps of about 0.1 °C, and on that grid the tuned rule avoids 10.5, 13.2 and 14.7% of on-commands and saves 5.7, 7.4 and 8.4% (at 23.6/22.4 °C in every city, which in Riyadh is the guard hysteresis, so on that grid the guard's shortfall F4 vanishes in Riyadh and the guard exceeds TUNED in Taif); under a per-window 0.03 K admission criterion it saves 9.7, 12.2 and 4.0%, 8.4 points less in Taif, where the guard then exceeds it (Supplementary Table S2). The headline TUNED figures are therefore an upper estimate of what a commissioned thermostat delivers. Any claim for a learned or model-based controller is measured against TUNED, not THERM.

Taif. The lower load does not leave less to win in percentage terms: TUNED saves 12.4% against a bracket of 16.7–21.0%, the widest of the three, with the lowest lower end and the highest upper end. The 24 h controller's 16.9% uses the simulator terminal; under V_N its equalised saving is 16.3%, below CLV's 16.5%.

Under the schedule announced for January 2027 (0.21 SAR/kWh) TUNED saves 11.6, 12.7 and 11.6%, and 18.6, 19.9 and 18.0% at the pre-2018 first-block price, while the on-commands it avoids are unchanged, because a higher energy price dilutes a saving whose mechanism is a fixed charge per on-command.

Table 3: Held-out results at the tariff in force, occupied band $[22, 24]\,^{\circ}\mathrm{C}$, $\Delta t = 15\,\mathrm{min}$, $c_{\mathrm{sw}} = 0.15\,\mathrm{SAR}$, single zone with one exposed face. s: pooled equalised saving (21) (%); $\Delta\tilde{T}$: pooled clipped-mean difference from THERM (K); on: change in on-commands against the equalised baseline (%). Fit status: ^{a}pooled 0.03 K; ^{b}pooled fallback 0.10 K; ^{o}oracle. Rows whose $\Delta\tilde{T}$ differs by more than about 0.02 K carry the comfort-offset caveat (about 1.3 points per 0.01 K, Section 10.4). MPC rows: 1,601 grid nodes, perfect (p) or persistence (s) forecast; †comfort matched by occupied-ceiling back-off b, because the λ device alone cannot reach the matched comfort level. "training rules": R1/R2 at 23.80/22.40 °C; "guard rules": each city's fitted guard thresholds. D1 rows at $\sigma = 0$. CLV: *raw* saving against unshifted THERM, the best attained schedule no warmer than THERM; its Rule C saving is paired with B_{CLV} in the text. Causal arms use the simulator's energy-only terminal; the CLV and bound rows use V_N (Section 5.5); B^* is given under both.

	Riyadh			Jeddah			Taif		
Arm	s	$\Delta\tilde{T}$	on	s	$\Delta\tilde{T}$	on	s	$\Delta\tilde{T}$	on
THERM (J, on/day)	1.118, 4.10			0.993, 3.73			0.614, 2.33		
TUNED	12.5^{a}	+0.011	−22.8	13.6^{a}	+0.029	−24.2	12.4^{a}	+0.005	−21.7
Guard hysteresis	5.7^{a}	+0.036	−10.5	7.6^{a}	+0.035	−13.6	11.0^{a}	+0.056	−19.3
RBRL, 5 seeds	5.7^{a}	+0.036	−10.5	7.6^{a}	+0.035	−13.6	8.6–11.0^{a}	—	—
D1 shield on latch	14.0^{a}	−0.009	−25.4	13.5^{a}	+0.021	−24.0	12.4^{a}	+0.005	−21.7
D1 on PPO, 5 seeds	12.8–14.1^{a}	—	—	11.2–13.5^{a}	—	—	8.0–12.0^{a}	—	—
MPC 0.25 h, p	-3.9^{a}	+0.058	+7.1	$1.6^{a\dagger}$	+0.023	−2.9	1.0^{a}	+0.034	−1.8
MPC 1 h, p	$10.1^{a\dagger}$	+0.024	−18.4	$10.1^{a\dagger}$	+0.001	−17.9	$6.8^{a\dagger}$	+0.009	−11.9
MPC 2 h, p	7.0^{b}	−0.099	−12.7	10.2^{a}	−0.062	−18.2	$9.2^{a\dagger}$	−0.003	−16.2
MPC 6 h, p	15.1^{a}	−0.021	−27.4	16.3^{a}	+0.022	−28.9	15.9^{a}	−0.005	−28.0
MPC 6 h, s	14.4^{a}	−0.004	−26.3	no admissible parameter			14.0^{a}	+0.003	−24.6
MPC 24 h, p	18.3^{a}	+0.014	−33.3	18.1^{a}	−0.011	−32.0	16.9^{a}	−0.011	−29.8
MPC 24 h, s	12.7^{b}	−0.076	−23.2	14.6^{a}	−0.037	−25.9	10.8^{b}	−0.086	−18.8
MPC 48 h, p	17.6^{a}	−0.018	−32.0	17.5^{a}	−0.012	−31.1	18.9^{b}	+0.089	−33.4
MPC 1 h, training rules	9.7^{b}	+0.115	−17.7	-3.6^{a}	−0.053	+6.5	-1.8^{b}	−0.023	+3.1
MPC 6 h, training rules	10.7^{a}	−0.028	−19.4	10.9^{a}	−0.006	−19.3	11.1^{b}	−0.016	−19.4
MPC 24 h, training rules	11.6^{b}	−0.061	−21.1	14.1^{a}	−0.003	−24.9	10.0^{b}	−0.064	−17.6
MPC 1 h, guard rules	2.5^{a}	+0.038	−4.5	1.5^{b}	+0.051	−2.8	8.0^{b}	+0.064	−14.2
MPC 6 h, guard rules	6.4^{a}	+0.025	−11.7	7.6^{a}	+0.024	−13.5	9.5^{b}	+0.076	−16.7
MPC 24 h, guard rules	6.7^{a}	+0.028	−12.3	4.4^{b}	+0.068	−7.7	10.5^{a}	+0.043	−18.6
CLV, raw^{o}	18.0	−0.036	−31.1	18.4	−0.018	−31.2	16.7	−0.025	−26.1
B^*, energy-only / V_N	≤ 20.0 / ≤ 20.0			≤ 20.1 / ≤ 19.9			— / ≤ 21.0		
CLV in rules^{o} / ruled bound	12.9 / ≤ 14.9			13.8 / ≤ 15.6			12.3 / ≤ 16.2		

Table 4: The three layers of every result, held-out windows, tariff in force. On: on-commands avoided against the comfort-equalised thermostat (%), the physical headline. Energy: electricity saved (%; negative means more), about unchanged for every arm. SAR: saving per room per cooling season, $184\,[c_{\mathrm{sw}}(\bar{f}_{\mathrm{eq}} - \bar{f}_{\mathrm{arm}}) + p(\bar{E}_{\mathrm{eq}} - \bar{E}_{\mathrm{arm}})]$ with $\bar{f}$ the mean on-commands per day and $\bar{E}$ the mean electricity per day over the seven held-out windows, $p = 0.18\,\mathrm{SAR/kWh}$, at $c_{\mathrm{sw}} = 0.025$–$0.114\,\mathrm{SAR}$ (the derived range) and, in brackets, at the 0.15 SAR computational reference, which overstates. Setback arms are scored on the occupancy timetable as scheduled, against THERM without setback equalised on the occupied steps.

	Riyadh			Jeddah			Taif		
Arm	On	Energy	SAR/season	On	Energy	SAR/season	On	Energy	SAR/season
TUNED	22.8	−0.3	4–19 (25)	24.2	+0.5	5–19 (25)	21.7	−0.5	2–10 (14)
Guard hysteresis	10.5	−0.4	2–8 (11)	13.6	0.0	2–10 (14)	19.3	0.0	2–9 (12)
D1 shield on latch	25.4	+0.7	5–22 (29)	24.0	−0.7	3–18 (24)	21.7	0.0	2–11 (14)
MPC 24 h, perfect	33.3	−0.4	6–28 (37)	32.0	+1.2	6–26 (34)	29.8	+0.8	4–15 (20)
MPC 24 h, persistence	23.2	−0.2	4–20 (26)	25.9	+0.7	5–21 (27)	18.8	−1.2	1–9 (12)
TUNED-SB (setback)	33.6	−0.2	6–28 (37)	35.0	+0.6	6–28 (36)	29.2	−0.8	3–14 (18)
MPC-SB (setback)	42.3	+0.6	9–37 (49)	39.7	+2.1	8–33 (42)	39.3	+3.3	6–21 (28)

Occupancy setback. With the dwelling empty from 08:00 to 14:00 on working days and an away limit of 30 °C, a tuned setback rule (TUNED-SB, with a fitted pre-arrival lead of 2, 1.25 and 1 h) saves 18.6, 20.1 and 17.0% ($\Delta\tilde{T} = +0.058, +0.027, +0.001\,\mathrm{K}$), and a 24 h receding-horizon controller with the timetable (MPC-SB) saves 24.0, 23.3 and 23.2% ($\Delta\tilde{T} = -0.030, -0.019, -0.076\,\mathrm{K}$), against setback bounds of 27.8, 27.6 and 28.8% and raw setback clairvoyant schedules of 25.3, 25.3 and 23.9%. Both are compared with THERM without setback, equalised on the occupied steps, so the

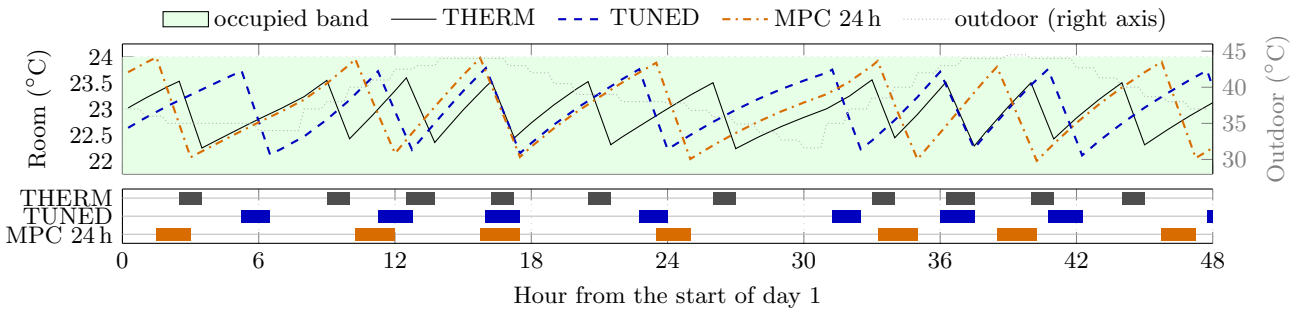


Fig. 3: A two-day operating trace: the hottest held-out Riyadh week, days 1–2, fixed before the traces were drawn (day 0 is warm-up and is not shown). Upper: room temperature at the end of each 15-min step under THERM, TUNED and the 24 h receding-horizon controller (perfect forecast), with the occupied band $[22, 24]\,^{\circ}$C shaded and the outdoor temperature (31.6–44.5 °C) on the right axis. Lower: cooling command (filled: on). THERM, TUNED and MPC issue 10, 8 and 7 on-commands over the two days and 34, 27 and 23 over the week. TUNED's switch points (23.71/22.36 °C) widen THERM's 1 °C deadband, so it cycles less; the controller uses the full band (22.00–23.99 °C).

saving includes letting the empty house float. The pre-arrival lead is not pre-cooling: it is the latest start that meets the 24 °C ceiling at arrival, and no cooling is banked beyond it.

Switching off at departure and on at arrival with no lead (MANUAL) is a reference, not an admitted arm: it leaves 21.8, 9.8 and 3.3 occupied degree-hours above the band after arrival, and its nominal saving (14.6, 13.0 and 13.7%) is bought with that over-temperature; THERM with the setback away limit coincides with it, because the away limit never binds (the empty zone peaks at 24.8–25.9 °C). The mechanism is on-commands avoided while the house is empty: TUNED-SB avoids 33.6, 35.0 and 29.2% of the equalised baseline's on-commands while energy moves by −0.2, +0.6 and −0.8%, and MPC-SB avoids 42.3, 39.7 and 39.3% with energy +0.6, +2.1 and +3.3% (vacation weeks: Supplementary Section S8). MPC-SB plans on the timetable it is given, and the model has no internal gains and a constant capacity, both of which flatter the recovery.

8.3 Where the Receding-Horizon Margin Comes From

The receding-horizon controller differs from TUNED in three ways at once: it has an exact model of the plant, it sees a forecast, and it plans over a horizon. The arms of Table 3 form a ladder that separates them: THERM, TUNED, the D1 shield (an exact one-step model with no horizon), MPC at 0.25, 1, 2, 6, 24 and 48 h with a perfect forecast, MPC with a persistence forecast, MPC confined to the rule set, E2-PPO, and finally CLV and the bound.

Against TUNED. With a perfect in-horizon forecast the 24 h controller avoids 33.3, 32.0 and 29.8% of on-commands against TUNED's 22.8, 24.2 and 21.7%, and its cost saving exceeds TUNED's in every held-out window of every city, by Hodges–Lehmann estimates of +6.0, +4.2 and +4.6 points (family F2, 95.3% intervals $[+3.8, +7.5]$, $[+3.8, +5.3]$ and $[+2.0, +8.0]$). With a persistence forecast the difference is −0.1, +0.9 and −1.9 points, with intervals that include zero and pooled differences under 3 points: not resolved (F3). With the one imperfect forecast tested, the margin over the tuned thermostat is therefore not resolved; forecasts between persistence and perfect were not tested. The persistence arms run 0.09, 0.07 and 0.09 K cooler than TUNED, in Riyadh and Taif at the 0.10 K fallback, which by the comfort-offset caveat penalises them; the back-off knob is tried only when no λ is admissible at 0.10 K, so it was not applied to these cells.

Horizon and model. The horizon sweep is descriptive (Table 3, with its back-off and fallback marks). Savings rise with the horizon up to 6 h, and differences of less than 3 points beyond it are not resolved; the Taif 48 h cell, fitted at the fallback and 0.09 K warmer than THERM, exceeds the clairvoyant lower end. The 0.25 h controller

saves at most 2% and costs more than the baseline in Riyadh (−3.9%). The one-step D1 shield on a latch saves 14.0, 13.5 and 12.4%, within 1.4 points of TUNED in every city, while the 24 h controller exceeds it by 4.3, 4.6 and 4.5 points: an exact one-step model reaches the tuned thermostat, and the margin beyond it needs the horizon and the forecast.

Forecast. At 6 h a persistence forecast costs 0.7 points in Riyadh and 1.9 in Taif, leaving 14.4 and 14.0%, 1.9 and 1.6 points above TUNED (not resolved); in Jeddah no multiplier or back-off is admissible at 6 h. At 24 h it costs 5.5, 3.4 and 6.2 points, partly at a cooler matched comfort (Table 3). Both modes plan on the plant's own model with the exact state, so neither measures model mismatch.

Thermal capacity. The single-node capacity is assumed, and the capacity effective at the cycle period is below the ISO 13786 value (Remark 1). In Riyadh, refitted at each capacity, TUNED saves −2.8, 12.5, 16.0 and 15.8% at 2,400, 4,800, 7,200 and 9,600 kJ/K, and the 24 h controller 12.2, 18.3, 19.8 and 16.2%, so its margin over TUNED is 15.0, 5.7, 3.8 and 0.4 points (Table 6); in Jeddah and Taif TUNED moves from −1.4 and −1.3% to 15.8 and 16.0% (16.7 and 17.7% at 7,200 kJ/K). The on-commands TUNED avoids rise monotonically with capacity in every city (Riyadh −4.0, 22.8, 34.6 and 39.8%). Its cost saving does not: the 7,200 and 9,600 kJ/K figures differ by less than 3 points in every city, and at 7,200 kJ/K TUNED sits 0.033, 0.040 and 0.121 K warmer than THERM against −0.019, +0.029 and −0.018 K at 9,600, so their order is not resolved. The planner's drop from 19.8 to 16.2% comes with a 0.044 K cooler matched mean and 2.5% more electricity at 9,600 kJ/K, while the on-commands it avoids barely move (42.4 and 40.7%). On the light plant tested (2,400 kJ/K) the commissioned thermostat does not beat a 1 °C deadband: each on-step overshoots by 0.75 K, TUNED's fitted deadband shrinks to 0.96 °C, and THERM's own fails (8) narrowly; THERM's self-score under Rule C was not computed there, so the small negative savings are not resolved from zero. On a heavy plant TUNED captures almost all of what the planner does. The ordering and its size therefore depend on the assumed capacity, and because the capacity acting over a cycle is probably below the reference (Remark 1), the reference results lean towards the end favourable to TUNED.

Robustness of the comfort device. The fitted multipliers of the 6 and 24 h perfect-forecast controllers lie at $\lambda/\lambda^* = 1.002$–$1.024$. In the following cells the fitted multiplier lies far above λ^*: $\lambda/\lambda^* = 1.226$ and 4.577 at 48 h in Riyadh and Jeddah, 13.094, 2.326 and 6.218 for the 24 h persistence controller, and 2.089 and 1.801 in rules at 24 h in Riyadh and Taif. Near λ^* the saving is sensitive to λ: neighbouring admissible multipliers move the Riyadh 24 h saving by 1.4–1.8 points. Runs at fixed $\lambda^* \cdot \{0.99, 1.01, 1.05\}$ are not comfort-matched and are given, with their $\Delta\bar{T}$, in Supplementary Table S2 only; grid refinement is reported in Section 5.5.

Inside the rules. Confined to the training rule set (23.80/22.40 °C) the 24 h controller saves 11.6, 14.1 and 10.0%, below relaxed bounds in those rules of 14.9, 15.6 and 16.2%. Confined to each city's *fitted* guard thresholds, the rules the guard hysteresis itself uses, it saves 6.7, 4.4 and 10.5% against 5.7, 7.6 and 11.0% for the guard: it does not beat the guard by a resolved margin in any city, and in Jeddah it falls 3.3 points short at a 0.03 K warmer mean, in a cell fitted at the 0.10 K fallback with $\lambda/\lambda^* = 0.60$. Inside the wider training rules, whose on threshold is 0.14–0.20 K higher, the 24 h controller exceeds the fitted guard by 5.9 and 6.4 points in Riyadh and Jeddah despite a 0.10 and 0.04 K cooler mean, which penalises it; this compares different rule sets. Replacing the guard's fixed margins by the D1 shield's one-step prediction lifts the latch to TUNED's level (14.0, 13.5 and 12.4%). Applied to the trained PPO policies in place of their rule layer, D1 gives 12.8–14.1 and 11.2–13.5% in Riyadh and Jeddah, the latch's level, but 8.0–12.0% in Taif, below the latch's 12.4%, and for three of the five seeds below their own rule-layer results (seed 42: 8.0 against 8.6%). Inside the guard's own fitted thresholds, then, neither lookahead nor learning adds a resolved margin, and the one-step shield sets the level (Remark 5).

8.4 Multi-Zone Dwellings

Whether the single-zone ordering survives in a dwelling of several rooms is tested in Riyadh for four layouts: the single zone, a 1 × 2 row, a 2 × 2 grid and a 1 × 4 row, each with the

glazed faces at the fixed azimuths of Table S-MZ and $\lambda_{ij} = 12.5\,\mathrm{W/K}$, with $35\,\mathrm{W/K}$ as a sensitivity. Every room has its own unit, so there is no shared capacity limit; rooms are coupled only through the shared walls. The arms are THERM, TUNED, the guard hysteresis and a multi-zone receding-horizon controller, all under the protocol of Section 7 applied to the dwelling's mean occupied temperature, with each room's mean and the spread between rooms reported beside it. The multi-zone controller is a non-cooperative distributed MPC with one-step-delayed plan exchange (Scattolini 2009) (Supplementary Section S7), with no optimality guarantee and no bound, because the joint dynamic programme is intractable (Section 9.4).

With rules alone, TUNED on the dwelling mean saves 12.5% in the single zone, 12.4% in the 1×2 row, 11.7% in the 2×2 grid and 13.6% in the 1×4 row (12.2, 12.9 and 13.3% at $\lambda_{ij} = 35\,\mathrm{W/K}$), and the guard hysteresis 5.7, 7.2, 8.3 and 7.7%; the mean spread of the room temperatures is 0.03–0.08 K. The multi-zone receding-horizon controller (24 h, perfect forecast) saves 16.0, 15.9 and 15.7% (15.8, 16.1 and 16.1% at $\lambda_{ij} = 35\,\mathrm{W/K}$), avoiding 28–32% of the equalised baseline's on-commands with electricity within ±0.7%; on the 1×1 instance it reproduces the single-zone 18.3% exactly (Table S-MZ). Its margin over TUNED is not tested for significance in the multi-zone layouts. Every room has the same band and timetable, so the rooms' mean temperatures differ by only 0.03–0.08 K and little heat crosses a wall: the study tests whether the single-zone ordering transfers to a coupled dwelling, not what coordination is worth.

9 State-of-the-Art Comparison and Ablation Study

This section compares the figures with published ones, ablates the rule layer and the learned component, gathers the sensitivities, and reports computational cost.

9.1 State-of-the-Art Comparison

Published RL-based HVAC studies differ in control type, building type, baseline and climate; Table 5 records the basis of each figure.

Xue et al. (2025), whose protocol is closest to this one (a rule-based comparator and a cost objective), report 6.7%. The figures here differ in three ways. Two lower them: the comparator is shifted to the arm's own mean temperature, and every figure is held out. One raises them: the saving is priced on on-commands at the 0.15 SAR reference, above the derived range, at which TUNED would save 3.8–11.0% in Riyadh. Electricity is about unchanged here, whereas the published figures are energy or cost savings. The larger published figures differ in scope, not degree, as the notes record.

9.2 Component Ablation and Sensitivity

The rule layer. Without R1 and R2, PPO trained on the same reward with a linear comfort penalty yields an admissible ceiling for one of five seeds in Riyadh, which then saves −3.5%, and for none of the five in Jeddah or Taif. That ceiling, 23.50 °C, is the lower edge of the search grid, so the null for the other nine runs may partly reflect the grid's range. This shows that this configuration fails to optimise without the rules; it does not show that a learner needs them.

The learned component. Under the rule layer, PPO on the cost-only reward reproduces the guard hysteresis in Riyadh and Jeddah: four of five seeds in each reproduce the guard's pooled saving to 10^{-5} (5.745 and 7.624%) and the fifth differs by less than 0.004 points. In Taif it does not: one seed equals the guard (11.0%) and the other four, at fitted ceilings of 23.84 or 23.86 °C, save 1.0–2.4 points less. Inside the rules hold is within 0.4, 1.4 and 2.3% of the $\lambda = 0$ dynamic-programming optimum in Riyadh, Jeddah and Taif (E1), so it is nearly optimal for the cost-only reward. E2, trained in Riyadh only with three seeds, a comfort term and a forecast, is not selected by Rule A and not comfort-matched; it stays within 0.16 on-commands per day and 2.6% of cost of the hold policy at the training rules. Because its multiplier sits near λ^*, where the λ term is flat in comfort, and hold is within 0.4% of the reward's optimum in Riyadh, E2 is consistent with, not evidence for, a learning null; its savings are not tabulated. Applying the exact one-step shield D1 in place of the five trained policies' rule layer gives 12.8–14.1 and 11.2–13.5% in Riyadh and Jeddah, the latch's

Table 5: State-of-the-art comparison. The "Basis" column is essential for interpreting saving figures.

Reference	Year	Method	Saving	Basis
(Azuatalam et al. 2020)	2020	RL, whole commercial building	22%	vs. handcrafted rule
(Du et al. 2021b)	2021	DDPG, multi-zone residential	15%	vs. DQN
(Biemann et al. 2021)	2021	SAC/TD3, continuous	>10%	vs. rule-based
(Wang et al. 2024)	2024	DRL, continuous VAV	37%[a]	vs. rule-based
(Solinas et al. 2024)	2024	Online RL, white-box	65%[b]	vs. no control
(Xue et al. 2025)	2025	GA-MADDPG hybrid	6.7%	vs. rule-based
This work	2026	TUNED / MPC-24 h / bracket, Riyadh	12.5 / 18.3 / 18.0–20.0%	vs. thermostat at equal mean, held out; on-commands priced at 0.15 SAR, electricity ≈0

[a] Continuous VAV modulation in a simulated office: different hardware and building type.
[b] Relative to a no-HVAC baseline; saving vs. thermostat is substantially lower.

level, and 8.0–12.0% in Taif, below it; the level is set by the shield, not by the policies, which on their tuning weeks already sit within 0.9% of the hold policy's training cost except one Taif seed (1.4%; Section 5.4).

The price of the guarantee. Compared with TUNED, which has no guard bands, the guard hysteresis costs 6.8 and 6.0 points pooled in Riyadh and Jeddah (Hodges–Lehmann 6.6 and 6.0), in every window (family F4); in Taif the difference, −1.4 points pooled, is not resolved, with the guard 0.05 K warmer than TUNED. On the 0.1 °C hardware grid the Riyadh difference vanishes, and under the per-window admission criterion the guard exceeds TUNED in Taif (Section 8.2). What the guards buy is a guarantee under the nominal plant parameters: nothing in TUNED's construction prevents an overtemperature or a cooling step below the band on unseen weather (it records the latter in Jeddah and Taif), whereas R1 and R2 exclude both by construction.

Sensitivities. Table 6 gathers the economic and physical parameters. The tariff and the charge per on-command move the percentage through the share of the baseline's cost that on-commands carry, as Remark 4 predicts. The comparator's deadband moves the saving most, from 40.2% at 0.5 °C to 5.7% at 1.25 °C, because it sets how often the baseline cycles: the reported figure is a statement about a pair, controller and comparator. TUNED was refitted at every charge; its switch points do not move for any positive charge, because the comfort limits bind before the charge does. The adopted $\delta_{\mathrm{db}} = 1\,^{\circ}\mathrm{C}$ is assumed, not measured. It sits near the upper end of the interval $[0.375, 1.25]\,^{\circ}\mathrm{C}$ admissible when the comparator is sampled at 15 min, a cap that a continuously sensing thermostat does not share, so an installed thermostat could have a wider deadband, against which TUNED saves less (5.7% against 1.25 °C). Measuring the installed stock's deadband is the field test most likely to confirm or overturn these results.

9.3 Actuation Step

Every figure above is tied to the 15-min actuation step. Table 7 and Fig. 4 repeat TUNED, the 24 h controller and the bound at 5 and 30 min, each arm refitted by Rule A. At 5 min the tuned thermostat avoids 43.1, 41.9 and 43.0% of on-commands, against 22.8, 24.2 and 21.7% at 15 min, with electricity within 1% of the equalised baseline's; its cost saving doubles to about 25%

Table 6: Sensitivity of the pooled equalised saving, Riyadh, seven held-out windows; descriptive, no tests. (a) Tariff schedule at $c_{\mathrm{sw}} = 0.15$ SAR. (b) Charge per on-command at the tariff in force; the derived range is 0.025–0.114 SAR. (c) THERM's deadband at a fixed $[22, 24]\,^{\circ}$C band; admissibility by (8); the comparator is sampled at 15 min. (d) Thermal capacity, every arm refitted by Rule A. Every TUNED and MPC cell is fitted at pooled 0.03 K; B^* under V_N. On-command share: $c_{\mathrm{sw}} f_s / J$ of THERM. Bold: reference setting. At 2,400 kJ/K THERM's 1 °C deadband fails (8) narrowly. TUNED is refitted at every charge; its switch points stay at 23.71/22.36 °C for every $c_{\mathrm{sw}} > 0$ because the comfort limits bind, and move to 23.54/22.39 °C at $c_{\mathrm{sw}} = 0$.

Setting	THERM on-command share	TUNED	MPC-24 h	B^*
(a) Tariff (SAR/kWh)				
0.05 (pre-2018)	81.5%	18.6	—	≤ 29.6
0.18 (in force)	55.0%	12.5	18.3	≤ 20.0
0.21 (2027)	51.2%	11.6	—	≤ 18.6
(b) Charge per on-command (SAR)				
0	0.0%	0.0	—	—
0.025	16.9%	3.8	5.7	≤ 6.3
0.05	29.0%	6.6	—	—
0.10	44.9%	10.2	—	—
0.114	48.2%	11.0	16.1	≤ 17.5
0.15	55.0%	12.5	18.3	≤ 20.0
0.20	62.0%	14.1	—	—
0.30	71.0%	16.2	—	—
0.60	83.0%	18.9	—	—
(c) THERM deadband (°C; THERM on-commands/day)				
0.50 (7.37)	—	40.2	—	—
0.75 (5.22)	—	23.4	—	—
1.00 (4.10)	—	12.5	18.3	—
1.25 (3.45)	—	5.7	—	—
1.50	inadmissible by (8)			
(d) Thermal capacity (kJ/K)				
2,400 (THERM 7.43/day)	—	−2.8	12.2	—
4,800 (4.10/day)	55.0%	12.5	18.3	≤ 20.0
7,200 (2.94/day)	—	16.0	19.8	≤ 20.0
9,600 (2.24/day)	—	15.8	16.2	—

and reaches 92% of the Riyadh bound. Its difference from the Riyadh 24 h controller is not resolved, and that controller ran on a coarser grid (801 nodes) at the fallback, 0.11 K cooler than TUNED. At 30 min TUNED no longer beats THERM: it issues 2.2–2.8% more on-commands and saves −1.2 to −1.5%, while the 24 h controller avoids 12.5, 21.0 and 20.4% and keeps 6.6–11.4%. There THERM's own deadband fails (8) narrowly and THERM leaks 0.025 occupied degree-hours in the first held-out Riyadh week, which favours THERM. The mechanism is the deadband of Section 8.2: the fitted deadband is 1.78–1.79 °C at 5 min and 0.94–0.99 °C, about THERM's own, at 30 min, as the per-step excursion allows. The step also changes the sampling of the comparator, a 1 °C hysteresis evaluated every Δt. The headline numbers therefore describe a 15-min supervisory step; they are not transferable to other control steps.

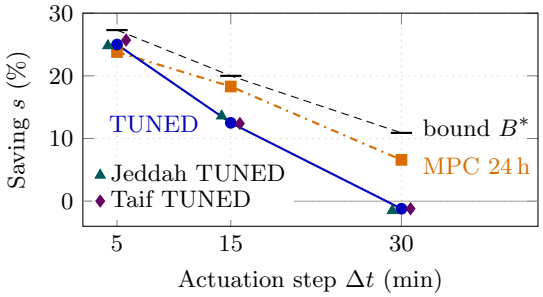


Fig. 4: Pooled equalised saving against actuation step, as in Table 7. Lines: Riyadh TUNED, 24 h receding-horizon controller (perfect forecast; at 5 min 801 nodes, 0.10 K fallback) and bound B^*; markers: Jeddah and Taif TUNED, offset horizontally. The gap between the tuned thermostat and the controller closes as the step shortens; at 30 min TUNED saves −1.2 to −1.5%.

Table 7: Saving s (%) against actuation step and capacity: TUNED, MPC-24 h (perfect forecast) and B^* (V_N, rounded up), every arm refitted by Rule A at pooled 0.03 K unless marked. Bold: headline step. —: not run; MPC and B^* at 5 min and at 7,200 kJ/K were run for Riyadh only, by design. At 30 min THERM's deadband fails (8) narrowly. [b]801 grid nodes (headline 1,601), pooled 0.10 K fallback, $\Delta\tilde{T} = -0.078$ K. [c]$\lambda/\lambda^* = 6.27$: for $\lambda \leq 0.21$ the controller holds the room about 0.2 K warmer than THERM on the tuning windows (pooled 23.12–23.14 against 22.91 °C, inadmissible), and at $\lambda \geq 0.29$ the mean jumps to 22.88 °C; the coarse step widens the sensitivity to λ of Section 8.3, the cheapest admissible λ is 1.51 (tuning cost 0.972 against 0.932 at the inadmissible 0.21), and the 6.6% is the protocol result, a conservative value. At 7,200 kJ/K TUNED is 0.033, 0.040 and 0.121 K warmer than THERM, so the comfort-offset caveat applies.

Δt, C_{th}	Arm	Riyadh	Jeddah	Taif
5 min, 4,800	TUNED	25.0	24.8	25.7
	MPC-24 h	23.8[b]	—	—
	B^*	$\leq$ 27.3	—	—
15 min, 4,800	TUNED	12.5	13.6	12.4
	MPC-24 h	18.3	18.1	16.9
	B^*	$\leq$ 20.0	$\leq$ 19.9	$\leq$ 21.0
30 min, 4,800	TUNED	−1.2	−1.5	−1.2
	MPC-24 h	6.6[c]	11.4	11.3
	B^*	$\leq$ 10.9	$\leq$ 13.9	$\leq$ 14.7
15 min, 7,200	TUNED	16.0	16.7	17.7
	MPC-24 h	19.8	—	—
	B^*	$\leq$ 20.0	—	—

9.4 Scalability

No claim is made about training cost, decision time or embedded deployability; the single-zone dynamic programme of the receding-horizon controller has 1,601 grid states and two actions per step.

What is settled here is structural. A clairvoyant programme over N_z coupled zones with G grid points per zone carries $G^{N_z}2^{N_z}$ states (the on-command charge makes each zone's cost depend on its previous action) and 2^{N_z} joint actions, so the backward pass touches $G^{N_z}4^{N_z}$ transitions per interval: 4.1×10^7 at two zones and $G = 1{,}601$, 2.6×10^{11} at three, 8.0×10^{60} at sixteen. That is why the bracket of Section 8.1 is a single-zone result and why the multi-zone study has no bound. Multi-zone results for a learned policy under the same tariff, with a larger unit, a $[23, 25]$ °C band and Jeddah profiles, are in (Faraj 2026) and are not directly comparable.

10 Discussion

This section interprets the city differences (Section 10.1), the robustness of the results (Section 10.2) and the factors that govern the saving (Section 10.3), and states the limitations (Section 10.4).

10.1 Climate and Headroom

The difference between the cities reflects load, diurnal swing and the baseline's cycling rate more than any algorithm. A low and variable load, as in Taif, does not mean little to win; it means that the switch points that win it may move from week to week, which a fixed pair cannot follow and a planner with a perfect forecast can (TUNED 12.4%, the 24 h controller 16.9%, with persistence 10.8%, against a bracket of 16.7–21.0%). Riyadh's 12.7 K diurnal swing lets the zone free-float further before the band forces the unit on; Jeddah's 8.4 K swing gives less. No schedule computed here pre-cools (Remark 2).

10.2 Robustness of the Reported Results

Five seeds per city were run and none was discarded; in Riyadh and Jeddah they reproduce the same policy, so their spread says nothing about robustness. The weather is varied through seven interleaved held-out weeks and three actual years per city (the mildest, median and hottest of fifteen by cooling degree-days at base 24 °C, from ERA5 reanalysis, which is not the station record behind TMYx, so the spread mixes year-to-year variation with dataset differences). TUNED's thresholds transferred unchanged from the typical year save 12.0–12.8, 13.2–13.4 and 10.5–12.0%, and retuned within each year 11.5–14.4, 12.3–14.8 and 13.7–14.4%; in Taif's two milder years TUNED free-floats below the band for 2.6–6.9 °C·h. Turning the Riyadh façade to north, east or west gives 13.3, 12.7 and 12.2% against 12.5% facing south (Supplementary Section S5). On calibration, no metered data exist for a dwelling of this description, so the model is an uncalibrated reference parameterised from the Saudi Building Code and rated equipment data; the ASHRAE Guideline 14 fit criteria cannot be computed, and no calibration claim is made. Absolute figures are not predictions for a real dwelling. Relative comparisons on

one model are not immune to a shared bias either: the margin of the receding-horizon controller over TUNED moves with the assumed capacity, and TUNED has no resolved saving at half the reference capacity (Section 8.3), so the parameters to which the ordering is sensitive are swept rather than assumed.

10.3 What Governs the Saving

In Riyadh the comparator's deadband moves the saving by 35 points across its admissible range, the controller class by 13 (guard 5.7, TUNED 12.5, 24 h controller 18.3%), the charge per on-command by 7 across the derived range and the tariff by 7 across the three epochs (Table 6). Only the controller class is algorithmic. The actuation step moves the saving further (Section 9.3).

The band acts mainly through the baseline. Narrowing it gives an optimising controller *less* room, but THERM, which switches 0.5 °C inside each edge, cycles more often on a narrower band, so the saving rises; the band was not swept under the protocol, so no figure is given. The narrower band is appropriate for this stock (Section 6.2), but any percentage reported for binary scheduling must be read with the band it was measured on.

What does move electricity is the comfort target. Along THERM's own equalisation curve on the held-out windows, inside the 24 °C ceiling, raising THERM's switching thresholds by 0.1 K (a 0.10–0.12 K warmer occupied mean) lowers its electricity by about 0.9, 1.4 and 2.3% in Riyadh, Jeddah and Taif, at the price of a warmer room; at matched comfort no switching policy changes it.

The tariff acts through the denominator. At matched mean temperature scheduling cannot change the energy delivered, only the number of on-commands, each worth c_{sw}; TUNED avoids the same 22.8% of the equalised baseline's on-commands in Riyadh under all three epochs, and the percentage falls as the bill grows (18.6, 12.5 and 11.6%). A percentage quoted against one tariff epoch is not comparable with one quoted against another. With a constant marginal price (Remark 3) and no time-of-use element, the store the dwelling holds (0.89 kWh electrical equivalent across the band) has nothing to be sold into; a time-varying price would change that.

Where, then, could a controller save money rather than on-commands? The receding-horizon controller and the clairvoyant bound already choose the whole temperature path inside the band at every step, coupled rooms included, and they do not pre-cool: with a single-speed unit, a constant COP and a flat price, cooling earlier only adds the energy the store leaks. Electricity at matched comfort becomes policy-dependent only when the cost of *delivered* cooling varies over the day, through a COP that rises at night (the 15–20% of Section 5.3), a time-of-use price, or a multi-speed or inverter unit whose part-load efficiency differs from full load. None of these is modelled here, so the conclusions hold for single-speed on/off units under a flat tariff; a controller that chooses speed and timing jointly across coupled rooms under those conditions, with setback in every room, is the natural next study.

The same distinction governs what the work offers the national targets. At matched comfort the controllers studied here change on-commands, not electricity, so their contribution to emission cuts is indirect, through fewer compressor cycles. Electricity, and with it emissions at a given generation mix, moves with the comfort target and with the outdoor air: lowering Riyadh's outdoor temperature by 1.5 and 2 K, the Green Riyadh aim, cuts THERM's cooling electricity on the held-out weeks by 10.5 and 13.7%, while TUNED's raw saving stays at 12.8–13.6%, so a cooler city and better switching add rather than compete. The shift is applied to an airport weather year with the plant and thresholds unchanged; it is a first-order sensitivity, not a model of the urban heat island.

10.4 Limitations

Supplementary Table S6 gives the direction in which each idealisation biases the reported percentage.

- *The saving is specific to six modelling choices.* The comparator's deadband, the actuation interval, the charge per on-command, the tariff, the assumed thermal capacity and the thermostat's setpoint resolution each move the headline by more than the resolution of the comparisons (Tables 6 and 7; Supplementary Tables S1 and S2), and the comfort band, which was not swept, acts on it through the baseline (Section 10.3); any figure quoted from this paper should be quoted with all of them. The deadband sweep is capped at 1.25 °C by

sampling the comparator at 15 min, and the comparator's 1 °C deadband is assumed, not measured.

- *Tied to the* 15*-min step.* The savings depend on the actuation step and the resolution at which the unit can be switched: TUNED saves about 25% at 5 min and −1.2 to −1.5% at 30 min (Section 9.3). The comparator is a 1 °C hysteresis evaluated at the actuation step; a thermostat that senses continuously is not represented, and the saving against it is not established. The headline numbers are not transferable to other control steps.
- *On-commands, not compressor starts.* The model issues supervisory on-commands to a lumped node; whether each is a compressor start depends on the unit's own thermostat, and the model's on-command rate is an artefact of that interface, not a field prediction.
- *Single-node thermal network and assumed capacity.* Air and fabric share one capacitance. Bacher and Madsen (2011), whose identification framework is followed, rejected the single-state model for the building they measured, because its residuals were not white, and selected a four-state grey-box model; a plant with a separate air node would cycle faster than this one. The capacity is assumed, and both results depend strongly on it (Section 8.3). ISO 13786's value is a 24 h quantity; the capacity acting at the period of THERM's cycle is probably lower, towards 2,400 kJ/K, where TUNED has no resolved saving, so the reference results lean towards the end favourable to TUNED.
- *Thermostat resolution.* TUNED's switch points are fitted on a 0.01 °C grid; restricted to the 0.1 °C steps of a real thermostat it saves 5.7, 7.4 and 8.4% instead of 12.5, 13.6 and 12.4%, so the commissioning gain depends on hardware that can resolve the thresholds. The guard and RBRL ceilings are fitted on the same fine grid. Comfort matching to hundredths of a kelvin and every switch point assume noise-free sensing for every arm, whereas typical thermistors err by a few tenths of a kelvin.
- *Perfect model and forecast.* The receding-horizon controller plans on the plant's own model with the exact state and, in its perfect mode, the true in-horizon weather; the persistence mode removes the second idealisation but not the first, and model mismatch is untested. Its terminal rollout uses persistence weather and the TUNED rule, so no closed-loop dominance over TUNED is guaranteed. The PPO agent's one-hour outdoor value is also exact.
- *Window-end terminal effect.* The terminal store is priced at re-cooling energy without an on-command charge, so an oracle that knows the window end can save up to one on-command per window, and in an earlier run under that terminal CLV ended its windows markedly warmer than THERM. The oracles therefore use V_N, which charges one on-command per kelvin left above the start temperature; that rate is a convention.
- *Plant data.* The 2 kW capacity and the COP of 2.985 are assumed, not taken from a rated unit. The COP lies below the SASO 2663:2018 minimum for split units at the T1 rating condition (about 3.46) though above its T3 minimum (about 2.43) (Saudi Standards, Metrology and Quality Organization 2018). A higher COP lowers the energy term and so raises the on-command share of cost and the percentage saving. Capacity and COP are constant, whereas a split unit loses capacity and efficiency at high outdoor temperatures; lower capacity acts in the hottest hours like a higher C_{th} in the actuation ratio. An outdoor-dependent COP (capacity held constant) is implemented in the repository but not reported; the effect of either derating on the comparison is not established.
- *No per-start energy and no internal gains.* An on-command costs money but no energy, so it cannot register as an energy saving; a transient start-up loss of 10–25 Wh would turn TUNED's energy difference into a saving of 0.33–0.80% in Riyadh, 0.37–0.90% in Jeddah and 0.34–0.83% in Taif. Occupancy, lighting and appliance gains, of order 150–300 W in an occupied room and so comparable with the modelled envelope load, are omitted, as is solar gain on opaque walls (an intermediate-floor room is modelled). They shorten free-float and lengthen on-phases, which is the mechanism the capacity and step sweeps act through; their effect on the headline savings is not established, and they flatter any setback's recovery.
- *Comfort offsets.* Rule C does not cancel a difference in comfort between arms: along the

guard family a 0.01 K warmer mean is worth about 1.3 points, so the 3-point resolution rule corresponds to about 0.02 K, below Rule A's 0.03 K tolerance. Arms whose $\Delta\tilde{T}$ differs by more than about 0.02 K, including those fitted at the 0.10 K fallback, are compared with that caveat. The admission rule itself matters: refitting under a per-window 0.03 K criterion moves TUNED by up to 8.4 points and reverses F4 in Taif (Supplementary Table S2); the receding-horizon controllers were not refitted under it, so the robustness of F2 and F3 to the admission rule is untested. Where the λ device cannot reach the matched comfort level, the receding-horizon controller is matched by backing off its occupied ceiling, and a negative saving there is a finding about the device, not a failure of the protocol.

- *Latent load.* The model is sensible-only. The latent axis was not tested at adequate fidelity, so for Jeddah, whose infiltration load is majority latent, this paper reports cost at matched dry-bulb comfort only, and not compliance with the ASHRAE 55 humidity criterion.
- *Scope of the learner.* PPO is tested only inside the rule layer and, in E2 (Riyadh, three seeds), with one comfort-aware reward; learners acting outside R1–R2 are untested; longer training is not run, but on the tuning weeks the trained policies already sit within 0.9% of the hold policy's training cost, except one Taif seed (1.4%), so it has little to gain. R1 and R2 impose no minimum dwell, so the complete-cycle structure is not enforced; a minimum-dwell rule set is required work.
- *One year, one zone type, simulation only.* The held-out weeks are interleaved within one TMYx year; the headline is a single zone with one exposed face; the multi-zone study is one city, independent units, the same band in every room and no bound, and a shared capacity-limited plant is not studied. Time-of-use or demand tariffs, minimum on- or off-times, a temperature-dependent COP and multi-speed or inverter units are not evaluated, and the conclusions do not extend to them; setback is evaluated in the single zone only. No result is reported on a community benchmark such as BOPTEST (Blum et al. 2021), and co-simulation and hardware-in-the-loop tests are required before any deployment.

11 Conclusions

Binary on/off cooling of Saudi residential split units was formulated as a constrained optimisation with hard comfort constraints, a charge per on-command and an energy cost, and every controller was selected and evaluated under one protocol in Riyadh, Jeddah and Taif. Four findings emerge.

First, commissioning a two-parameter thermostat widens its deadband from the assumed 1 °C to 1.32–1.36 °C and avoids 22–24% of on-commands in every held-out window, with electricity within ±0.5%; priced at the derived 0.025–0.114 SAR per on-command this is worth 2–19 SAR per room per cooling season. The gain is a statement about the comparator as much as about tuning: against a 1.25 °C deadband the tuned thermostat saves 5.7% in Riyadh instead of 12.5%, on a 0.1 °C set-point grid it avoids 11–15% of on-commands, and at its own mean temperature it attains 62–65% of the provable bound.

Second, the comparison depends strongly on the plant and the step. A 24 h receding-horizon controller avoids 30–33% of on-commands and adds 4–6 points of cost saving with a perfect in-horizon forecast; with a persistence forecast, the only imperfect one tested, its margin is not resolved. In Riyadh the tuned thermostat has no resolved saving at half the reference capacity and saves 15.8% at twice it, while the controller's margin falls from 15.0 to 0.4 points; because the capacity acting over a cycle is probably below the reference, the reference results lean towards the end favourable to the tuned thermostat. At a 5-min step it saves about 25% and at 30 min −1.2 to −1.5%, so every figure here holds for a 15-min step only.

Third, PPO trained under hard rules reproduces the rules' hold policy in Riyadh and Jeddah and falls below it in Taif, and a comfort-aware variant, tested in Riyadh only, stays close to hold. Inside the guard's own fitted thresholds neither lookahead nor learning adds a resolved margin; an exact one-step shield in place of the rules' fixed margins reaches the tuned thermostat's level on a latch, and inside the wider training rules

a 24 h controller gains about 6 points over the guard in Riyadh and Jeddah. The guarantee the guards give costs 6.8 and 6.0 points in Riyadh and Jeddah.

Fourth, at matched mean temperature every arm uses about the same electricity, by an identity of the model, so the cost difference is a charge on avoided on-commands, assumed to stand for wear; electricity follows the comfort target, each 0.1 K warmer mean lowering THERM's by 0.9–2.3%, and with the outdoor air, which 1.5–2 K cooler (the Green Riyadh aim) lowers it by 10.5–13.7% in Riyadh. The money saved depends on the six modelling choices of Section 10.4 and on the comfort band, which was not swept. Occupancy setback raises the on-commands the tuned rule avoids to 29–35% and its cost saving to 17–20%. In Riyadh dwellings of two and four rooms with independent units the tuned thermostat saves 11.7–13.6%, against 12.5% in the single zone, differences that are not resolved, and a non-cooperative distributed receding-horizon controller saves 15.7–16.1%, against 18.3% in the single zone.

For this model instance three recommendations follow. (i) Judge any learner or planner against a thermostat whose deadband has been widened to what the band and the actuation step admit, not against the factory deadband. In practice such commissioning assumes a supervisory 15-min control layer and thresholds resolved to 0.01 °C (on a 0.1 °C grid the saving is 5.7–8.4%), and no commissioning or hardware cost is set here against the 2–19 SAR per season. (ii) Check both sides of the actuation condition (8) and derive guard bands from the per-interval excursion in each direction before trusting any reported saving. (iii) Report a saving with the band, the comparator's deadband, the charge per on-command and the on-commands avoided, and do not present it as a bill saving. Every figure is simulated, and field validation, with the installed stock's deadband measured, is required before any deployment claim.

Funding. This work was funded by the Deanship of Graduate Studies and Scientific Research, Taif University.

Author contributions. **Hamzah Faraj:** Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Writing – original draft, Writing – review & editing, Visualization.

Data availability. *What is released.* The code is available at the project repository on GitHub, archived on Zenodo at doi:10.5281/zenodo.24199676 (release v4.0-round4); the trained PPO models (29 MB) are not in the archive, are available from the author on request, and can be retrained from the code with the stated seeds. It comprises the thermal model, every controller and oracle, the protocol module that applies Rules A–C to every arm, the statistics, and the scripts that generate every table from the results files, each of which records the code version that produced it. A gate of 188 numerical invariants in nine classes (physical sanity, bound consistency, accounting identities, parameter provenance, monotonicity, decomposition closure, feasibility, the deadband actuation condition, and the protocol checks on every results file) runs before any table is produced and exits non-zero on any failure; the feasibility class retains some checks written for the band of an earlier version of the model. The time-step and capacity comfort records of THERM in Supplementary Table S1 come from a separate THERM run on the first held-out Riyadh week. Code paths that the paper does not report (an outdoor-dependent COP, a behaviour-cloning variant, a shared-plant multi-zone arbitration) are in the repository and marked as unreported.

What is not available. An exploratory pipeline that preceded this study is archived in the repository (`archive/withdrawn-pipeline`) for transparency only; it does not reproduce, and is not the source of, any number or figure reported here.

Weather data. All three records are TMYx files from climate.onebuilding.org: Riyadh (WMO 404380, 2009–2022), Jeddah (WMO 410240, 2009–2023) and Taif (WMO 410360, 2009–2023) (Lawrie and Crawley 2022). The source reserves its rights, so the EPW files are not redistributed; `fetch_weather.py` names each station file and verifies the downloaded copy by SHA-256. The actual-year series of the inter-annual study are ERA5 reanalysis (Hersbach et al. 2020) retrieved through the Open-Meteo archive API, redistributed in the repository under CC BY 4.0 with attribution (weather data by Open-Meteo.com; contains

modified Copernicus Climate Change Service information).

Declarations

Ethics approval and consent to participate. This research does not involve human participants, animal subjects or sensitive data. Ethics approval and consent to participate are not applicable; consent for publication is not applicable.

Competing interests. The author declares no competing interests.

Use of generative AI. During the preparation of this work the author used Claude (Anthropic) for language editing, code assistance and simulated peer review of the manuscript. All analyses, results and conclusions were checked by the author, who reviewed and edited the content as needed and takes full responsibility for the publication.

References

Al-Homoud MS, Krarti M (2021) Energy efficiency of residential buildings in the Kingdom of Saudi Arabia: Review of status and future roadmap. Journal of Building Engineering 36:102143. https://doi.org/10.1016/j.jobe.2020.102143

Al Sayed K, Boodi A, Sadeghian Broujeny R, et al (2024) Reinforcement learning for HVAC control in intelligent buildings: a technical and conceptual review. Journal of Building Engineering 95:110085. https://doi.org/10.1016/j.jobe.2024.110085

Alayed E, O'Hegarty R, Kinnane O (2021) Thermal envelope analysis for new code compliance of Saudi Arabian dwellings. Energy and Buildings 243:110997. https://doi.org/10.1016/j.enbuild.2021.110997

Alotaibi AM, Makhdoom TK, Alquaity ABS (2024) Pathways toward improving the energy efficiency of residential air-conditioning systems in Saudi Arabia. Journal of Solar Energy Engineering 146(5):051010. https://doi.org/10.1115/1.4065973

Alshiekh M, Bloem R, Ehlers R, et al (2018) Safe reinforcement learning via shielding. In: Proceedings of the AAAI Conference on Artificial Intelligence, pp 2669–2678, https://doi.org/10.1609/aaai.v32i1.11797

Argaam (2024) Cabinet OKs amending article 9 of Saudi Building Code application law. https://www.argaam.com/en/article/articledetail/id/1761923, amendment to the Law on Implementation of the Saudi Building Code (Royal Decree M/43, 2017): the occupancy certificate is a prerequisite for full electricity activation. Accessed 11 November 2026

ASHRAE (2023) ANSI/ASHRAE Standard 55-2023: Thermal Environmental Conditions for Human Occupancy. American Society of Heating, Refrigerating and Air-Conditioning Engineers, Peachtree Corners, GA, URL https://www.ashrae.org/technical-resources/bookstore/standard-55-thermal-environmental-conditions-for-human-occupancy

Azuatalam D, Lee WL, de Nijs F, et al (2020) Reinforcement learning for whole-building HVAC control and demand response. Energy and AI 2:100020. https://doi.org/10.1016/j.egyai.2020.100020

Bacher P, Madsen H (2011) Identifying suitable models for the heat dynamics of buildings. Energy and Buildings 43(7):1511–1522. https://doi.org/10.1016/j.enbuild.2011.02.005

Bandarra Filho EP, Martins G, Riaz MB, et al (2025) Optimization of air conditioning mechanical ventilation using simulated annealing for enhanced energy efficiency and cost reduction. Scientific Reports 15:23429. https://doi.org/10.1038/s41598-025-07640-z

Belaïd F, Fikru MG (2025) Resolving the energy efficiency paradox: leveraging benefits for Saudi Arabia's building sector. Humanities and Social Sciences Communications 12:1573. https://doi.org/10.1057/s41599-025-05078-2

Belaïd F, Mikayilov JI (2024) Closing the efficiency gap: insights into curbing the direct rebound effect of residential electricity consumption in Saudi Arabia. Energy Economics

135:107647. https://doi.org/10.1016/j.eneco.2024.107647

Bertsekas DP (2020) Rollout, Policy Iteration, and Distributed Reinforcement Learning. Athena Scientific, Belmont, MA

Biemann M, Scheller F, Liu X, et al (2021) Experimental evaluation of model-free reinforcement learning algorithms for continuous HVAC control. Applied Energy 298:117164. https://doi.org/10.1016/j.apenergy.2021.117164

Bienvenido-Huertas D, Sánchez-García D, Tejedor B, et al (2024) Energy savings in buildings applying ASHRAE 55 and regional adaptive thermal comfort models. Urban Climate 55:101892. https://doi.org/10.1016/j.uclim.2024.101892

Blum D, Arroyo J, Huang S, et al (2021) Building optimization testing framework (BOPTEST) for simulation-based benchmarking of control strategies in buildings. Journal of Building Performance Simulation 14(5):586–610. https://doi.org/10.1080/19401493.2021.1986574

Di Natale L, Svetozarevic B, Heer P, et al (2023) Towards scalable physically consistent neural networks: an application to data-driven multi-zone thermal building models. Applied Energy 340:121071. https://doi.org/10.1016/j.apenergy.2023.121071

Ding X, Cerpa A, Du W (2025) Multi-zone HVAC control with model-based deep reinforcement learning. IEEE Transactions on Automation Science and Engineering 22:4408–4426. https://doi.org/10.1109/TASE.2024.3410951

Dinh HT, Kim D, Kim D (2022) MILP-based optimal day-ahead scheduling for a system-centric community energy management system supporting different types of homes and energy trading. Scientific Reports 12:18305. https://doi.org/10.1038/s41598-022-22293-y

Drgoňa J, Arroyo J, Cupeiro Figueroa I, et al (2020) All you need to know about model predictive control for buildings. Annual Reviews in Control 50:190–232. https://doi.org/10.1016/j.arcontrol.2020.09.001

Drgoňa J, Tuor AR, Chandan V, et al (2021) Physics-constrained deep learning of multi-zone building thermal dynamics. Energy and Buildings 243:110992. https://doi.org/10.1016/j.enbuild.2021.110992

Du Y, Li F, Munk J, et al (2021a) Multi-task deep reinforcement learning for intelligent multi-zone residential HVAC control. Electric Power Systems Research 192:106959. https://doi.org/10.1016/j.epsr.2020.106959

Du Y, Zandi H, Kotevska O, et al (2021b) Intelligent multi-zone residential HVAC control strategy based on deep reinforcement learning. Applied Energy 281:116117. https://doi.org/10.1016/j.apenergy.2020.116117

El Makroum R, Khallaayoun A, Lghoul R, et al (2023) Home energy management system based on genetic algorithm for load scheduling: a case study based on real life consumption data. Energies 16(6):2698. https://doi.org/10.3390/en16062698

Faddel S, Sun Q, Tian G, et al (2023) Scheduling of the HVAC system in a real commercial building considering equipment cycling and rebound effects. Frontiers in Energy Research 11:1283369. https://doi.org/10.3389/fenrg.2023.1283369

Faraj H (2026) Cost-effective tariff-aware HVAC coordination in Saudi residential buildings using physics-informed deep reinforcement learning. Journal of King Saud University – Engineering Sciences 38:58. https://doi.org/10.1007/s44444-026-00131-0

Haddad S, Zhang W, Paolini R, et al (2024) Quantifying the energy impact of heat mitigation technologies at the urban scale. Nature Cities 1:62–72. https://doi.org/10.1038/s44284-023-00005-5

Hersbach H, Bell B, Berrisford P, et al (2020) The ERA5 global reanalysis. Quarterly Journal of the Royal Meteorological Society 146(730):1999–2049. https://doi.org/10.1002/qj.3803

Hou F, Cheng JCP, Kwok HHL, et al (2024) Multi-source transfer learning method for enhancing the deployment of deep reinforcement learning in multi-zone building HVAC control. Energy and Buildings 322:114696. https://doi.org/10.1016/j.enbuild.2024.114696

Huang S, Ontañón S (2022) A closer look at invalid action masking in policy gradient algorithms. In: Proceedings of the International FLAIRS Conference, https://doi.org/10.32473/flairs.v35i.130584

Krarti M (2020) Evaluation of occupancy-based temperature controls on energy performance of KSA residential buildings. Energy and Buildings 220:110047. https://doi.org/10.1016/j.enbuild.2020.110047

Krarti M, Aldubyan M, Williams E (2020) Residential building stock model for evaluating energy retrofit programs in Saudi Arabia. Energy 195:116980. https://doi.org/10.1016/j.energy.2020.116980

Lamsal P, Bajracharya SB, Rijal HB (2023) A review on adaptive thermal comfort of office building for energy-saving building design. Energies 16(3):1524. https://doi.org/10.3390/en16031524

Lawrie LK, Crawley DB (2022) Development of global typical meteorological years (TMYx). https://climate.onebuilding.org, tMYx EPW files used: Riyadh WMO 404380 (2009–2022), Jeddah WMO 410240 (2009–2023), Taif WMO 410360 (2009–2023). Accessed 2026

Liu X, Wu Y, Wu H (2024) Enhancing HVAC energy management through multi-zone occupant-centric approach: a multi-agent deep reinforcement learning solution. Energy and Buildings 303:113770. https://doi.org/10.1016/j.enbuild.2023.113770

Manjavacas A, Campoy-Nieves A, Jiménez-Raboso J, et al (2024) An experimental evaluation of deep reinforcement learning algorithms for HVAC control. Artificial Intelligence Review 57(7):173. https://doi.org/10.1007/s10462-024-10819-x

Mousa MAA, Schewe S, Wojtczak D (2016) Optimal control for simple linear hybrid systems. In: Proceedings of the 23rd International Symposium on Temporal Representation and Reasoning (TIME). IEEE, pp 12–20, https://doi.org/10.1109/TIME.2016.9

Mousa MAA, Schewe S, Wojtczak D (2017) Optimal control for multi-mode systems with discrete costs. In: Formal Modeling and Analysis of Timed Systems (FORMATS 2017), Lecture Notes in Computer Science, vol 10419. Springer, pp 77–96, https://doi.org/10.1007/978-3-319-65765-3_5

Nahiduzzaman KM, Said Abdallah A, Moradzadeh A, et al (2023) Impacts of tariffs on energy conscious behavior with respect to household attributes in Saudi Arabia. Energies 16(3):1458. https://doi.org/10.3390/en16031458

Olawale MA, Al-Homoud MS, Abdou AA, et al (2025) A comprehensive assessment of the impact of the Saudi Building Code (SBC) on energy performance of residential buildings in Saudi Arabia. Energy and Buildings 349:116524. https://doi.org/10.1016/j.enbuild.2025.116524

Pang X, Liu M (2011) Prevention of compressor short cycling in direct-expansion (DX) rooftop units, Part 1: Theoretical analysis and simulation. ASHRAE Transactions 117(2):666–676

Pérez-Lombard L, Ortiz J, Pout C (2008) A review on buildings energy consumption information. Energy and Buildings 40(3):394–398. https://doi.org/10.1016/j.enbuild.2007.03.007

Raffin A, Hill A, Gleave A, et al (2021) Stable-Baselines3: Reliable reinforcement learning implementations. Journal of Machine Learning Research 22(268):1–8. URL http://jmlr.org/papers/v22/20-1364.html

Risbeck MJ, Maravelias CT, Rawlings JB, et al (2017) A mixed-integer linear programming model for real-time cost optimization of building heating, ventilation, and air conditioning equipment. Energy and Buildings 142:220–235. https://doi.org/10.1016/j.enbuild.2017.02.053

Rodrigues L, Cherian BA, Tokbolat S (2025) Reducing cooling energy demand in Saudi Arabian residential buildings using passive design approaches. Buildings 15(11):1895. https://doi.org/10.3390/buildings15111895

Saudi Energy Company (2026) Revised electricity, water and wastewater tariffs effective 1 january 2027. https://ajel.sa/economy/akkj28rot, announced 25 October 2026. Residential consumption tariff repriced to 0.21 SAR/kWh up to 6,000 kWh/month and 0.33 SAR/kWh above, with a fixed monthly charge of 70–100 SAR graded by breaker capacity replacing the former meter reading, maintenance and billing fee. Accessed 11 November 2026

Saudi Gazette (2025) Revised residential and commercial electricity tariffs to remain same. https://saudigazette.com.sa/article/652853/SAUDI-ARABIA/Revised-residential-and-commercial-electricity-tariffs-to-remain-same, residential: 0.18 SAR/kWh up to 6,000 kWh/month, 0.30 SAR/kWh above. Accessed 11 November 2026

Saudi Green Initiative (2021) About the Saudi Green Initiative: fact sheet. https://sgi.gov.sa/media/v4pllu0w/about_saudi_green_initiative_factsheet_en.pdf, accessed 13 November 2026

Saudi Standards, Metrology and Quality Organization (2018) SASO 2663:2018 – minimum energy performance requirements and labelling for small air conditioners. Summary in the IEA Policies Database, URL https://www.iea.org/policies/8944-saso-26632018-meps-small-air-conditioners, accessed 12 November 2026

Saudipedia (2019) Green Riyadh project. https://saudipedia.com/en/green-riyadh-project, launched 19 March 2019; accessed 13 November 2026

Scattolini R (2009) Architectures for distributed and hierarchical model predictive control – a review. Journal of Process Control 19(5):723–731. https://doi.org/10.1016/j.jprocont.2009.02.003

Shi H, Liu J, Chen Q (2022) An RC-network approach for HVAC precooling optimization in buildings. IEEE Transactions on Sustainable Computing 7(3):512–526. https://doi.org/10.1109/TSUSC.2019.2943491

Solinas FM, Macii A, Patti E, et al (2024) An online reinforcement learning approach for HVAC control. Expert Systems with Applications 238:121749. https://doi.org/10.1016/j.eswa.2023.121749

Wang H, Chen X, Vital N, et al (2024) Energy optimization for HVAC systems in multi-VAV open offices: a deep reinforcement learning approach. Applied Energy 356:122354. https://doi.org/10.1016/j.apenergy.2023.122354

Wang J, Jiang Y, Tang CY, et al (2022) Development and validation of a second-order thermal network model for residential buildings. Applied Energy 306:118124. https://doi.org/10.1016/j.apenergy.2021.118124

Wang Z, Hong T (2020) Reinforcement learning for building controls: The opportunities and challenges. Applied Energy 269:115036. https://doi.org/10.1016/j.apenergy.2020.115036

Xu S, Fu Y, Wang Y, et al (2025) Efficient and assured reinforcement learning-based building HVAC control with heterogeneous expert-guided training. Scientific Reports 15:7677. https://doi.org/10.1038/s41598-025-91326-z

Xue W, Jia N, Zhao M (2025) Multi-agent deep reinforcement learning based HVAC control for multi-zone buildings considering zone-energy-allocation optimization. Energy and Buildings 329:115241. https://doi.org/10.1016/j.enbuild.2024.115241

Yu L, Sun Y, Xu Z, et al (2021) Multi-agent deep reinforcement learning for HVAC control in commercial buildings. IEEE Transactions on Smart Grid 12(1):407–419. https://doi.org/10.1109/TSG.2020.3011739